

# Week 6

# System Call Programming

5 November 2018

CS 35L Lab 4

Jeremy Rotman

# Announcements

- Assignment #5 is due Saturday by 11:55pm
- Assignment #10 Presentations
  - ◆ **Email me to tell me what story you are choosing**
  - ◆ [Here is the link to see what stories people have signed up for already](#)
  - ◆ [Here is the link to sign up to present](#)
- Remember to change your clocks after the extra hour we received this weekend
- Tomorrow is Election Day

# Outline

- Dual-Mode Operation
- System Calls
- Buffered vs. Unbuffered I/O
- Lab 5

Questions?

# Processor Modes

→ The CPU (in Linux) has two distinct modes of operation

## ◆ Kernel mode

- Unrestricted access
- Can execute any instruction, and reference any memory address
- Assumes it is running trusted software

## ◆ User mode

- Non-privileged access
- Cannot directly access hardware
- Must use a system call to perform privileged instructions

# Why Dual-Mode Operation?


- I/O protection
  - ◆ Input/output cannot be directly controlled by user-created code
- Memory protection
  - ◆ User mode only has access to a set partition of memory
  - ◆ The user is not allowed to access memory addresses that define these bounds, or other addresses outside their partition
- CPU Protection
  - ◆ User mode is not allowed to change things related to the OS's scheduler or timer

# Trusted Software

- What is the trusted software that can be run in kernel mode?

# Trusted Software

- What is the trusted software that can be run in kernel mode?
  - ◆ Software in the kernel space
    - Cannot be changed from the outside
    - Implements protection mechanisms
- System call interface bridges the gap between User Mode and Kernel Mode
  - ◆ User processes can execute privileged operations through the interface






# System Calls

- Used by user-level processes to request service from kernel
- Changes CPU's mode from User to Kernel
- Part of the Kernel of the OS
- Verifies User should be allowed to do operation, then handles operation
- The only way that a user program can perform privileged operations

# System Calls

- When a system call is made
- ◆ The program being executed is interrupted
  - ◆ Control is passed to the kernel
  - ◆ If the operation is valid
    - Kernel performs it
  - ◆ Else
    - Throw an exception



# System Calls

## → Overhead

- ◆ System calls include a lot of overhead
- ◆ Many things that must be done
  - Process interrupted and computer saves its state
  - OS takes control of CPU and verifies validity of operation
  - OS performs requested action
  - OS restores saved context, switches to user mode
  - OS gives control of the CPU back to user process

# Example System Calls

- `ssize_t read(int fildev, void* buf, size_t nbyte)`
  - ◆ *fildev*: file descriptor
  - ◆ *buf*: address of buffer to read into
  - ◆ *nbyte*: maximum number of bytes to read
- `ssize_t write(int fildev, void* buf, size_t nbyte)`
  - ◆ *fildev*: file descriptor
  - ◆ *buf*: address of buffer to read into
  - ◆ *nbyte*: maximum number of bytes to write
- `int open(const char* pathname, int flags, mode_t mode)`
- `int close(int fildev)`
- `fd: 0 = stdin; 1 = stdout; 2 = stderr;`

# Example System Calls

- `void exit(int status)`
  - ◆ Terminates process with *status*
- `pid_t fork(void)`
  - ◆ Creates child process
- `pid_t getpid(void)`
  - ◆ Returns the process ID of the calling process
- `int dup(int fildes)`
  - ◆ Duplicates a file descriptor, *fildes*
- `int fstat(int fildes, struct stat* buf)`
  - ◆ Return information about the file with file descriptor *fildes*, into *buf*

# C Library Functions

- There are functions that are part of the standard C library that do the same thing
  - ◆ getchar & putchar are similar to read & write (standard I/O)
  - ◆ fopen & fclose are similar to open & close (file I/O)
- These C functions then make system calls of their own
- What is the advantage?

# C Library Functions

- There are functions that are part of the standard C library that do the same thing
  - ◆ getchar & putchar are similar to read & write (standard I/O)
  - ◆ fopen & fclose are similar to open & close (file I/O)
- These C functions then make system calls of their own
- What is the advantage?
  - ◆ Library functions make fewer system calls
  - ◆ This reduces the amount of overhead
  - ◆ Efficiency :D

# Unbuffered vs. Buffered I/O

## → Unbuffered

- ◆ Every byte is read/written by the kernel through system call

## → Buffered

- ◆ Collect as many bytes as possible (in a buffer) and read/write them all in one system call

## → Buffered I/O reduces overhead because you don't have to switch mode for every byte

- ◆ Thought exercise: when might you want to use unbuffered I/O?
- ◆ Also, you must still always be prepared for buffer overflow attacks



# Lab 5

→ Write C code to transliterate bytes

◆ 2 arguments:

- *from*: the bytes to transliterate
- *to*: the bytes to transliterate them into

◆ Works basically the same way as `tr`

- `tr 'abcd' 'nopq' < input_file`
  - Replace 'a' with 'n', 'b' with 'o', etc.

→ You will write 2 programs (one buffered, one not)

◆ `tr2b`

- Uses `getchar` and `putchar` to read from `stdin` and write to `stdout`

◆ `tr2u`

- Uses `read` and `write` with the `nbyte` argument as 1

# Lab 5

- The programs will need to be tested on a very large file
  - ◆ At least 5,000,000 bytes
  - ◆ `head --bytes=num /dev/urandom > output.txt`
- **time** [options] *command* [arguments...]
  - ◆ Output:
    - `real 0m0.145s`
      - Elapsed real time (in [hours:]minutes:seconds)
    - `user 0m0.001s`
      - Total number of CPU-seconds that process spent in user mode
    - `sys 0m0.003s`
      - Total number of CPU-seconds that process spent in kernel mode

# Lab 5

- **strace** -o *strace\_output command* [arguments...]
- ◆ Intercepts and prints out system calls to stderr or to an output file
  - ◆ Every line will include
    - system call name
    - arguments in parentheses
    - return value

Questions?