

# CS35L – Fall 2018

Slide set:	6.1
Slide topics:	Multithreaded performance
Assignment:	6

# Parallelism

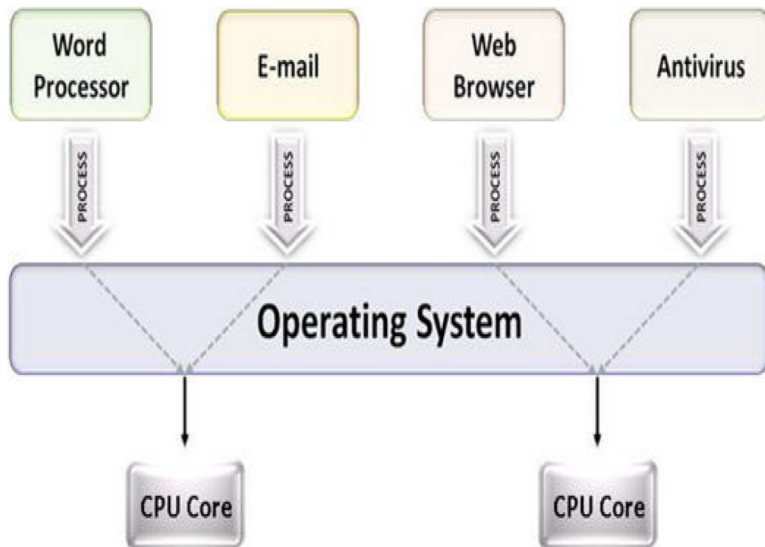
- Executing several computations simultaneously to gain performance
- Different forms of parallelism
  - **Multitasking**
    - Several processes are scheduled alternately or possibly simultaneously on a multiprocessing system
  - **Multithreading**
    - Same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system

# What is a thread?

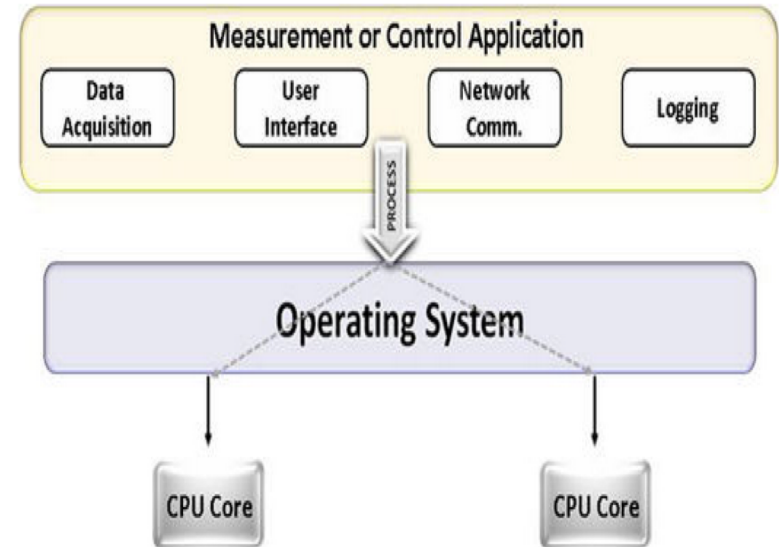
- A flow of instructions, path of execution within a process
- The smallest unit of processing scheduled by OS
- A process consists of at least one thread
- Multiple threads can be run on:
  - **A uniprocessor (time-sharing)**
    - Processor switches between different threads
    - Parallelism is an illusion
  - **A multiprocessor**
    - Multiple processors or cores run the threads at the same time
    - True parallelism

# Multitasking vs. Multithreading

## Multitasking




## Multithreading




Multitasking is sharing of computing resources(CPU, memory, devices, etc.) among processes


Multithreading is sharing of computing resources among threads of a single process.





one process  
one thread




one process  
multiple threads




multiple processes  
one thread per process



multiple processes  
multiple threads per process



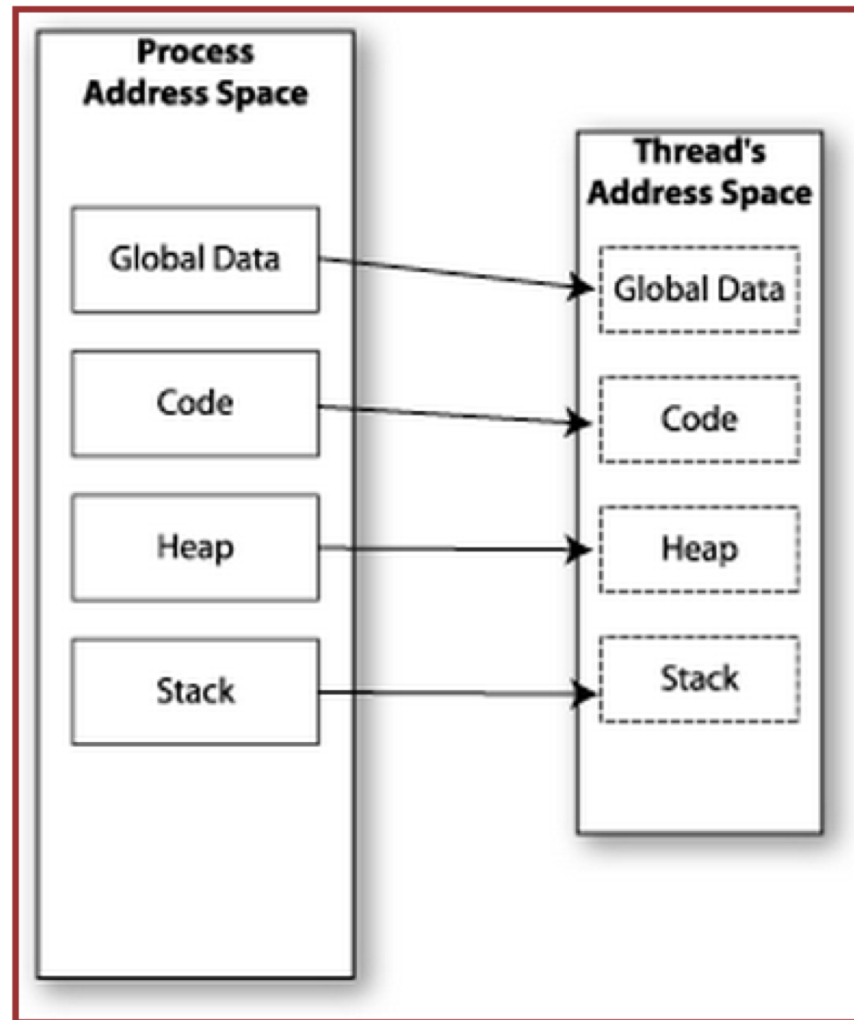

Multiple threads sharing a single CPU



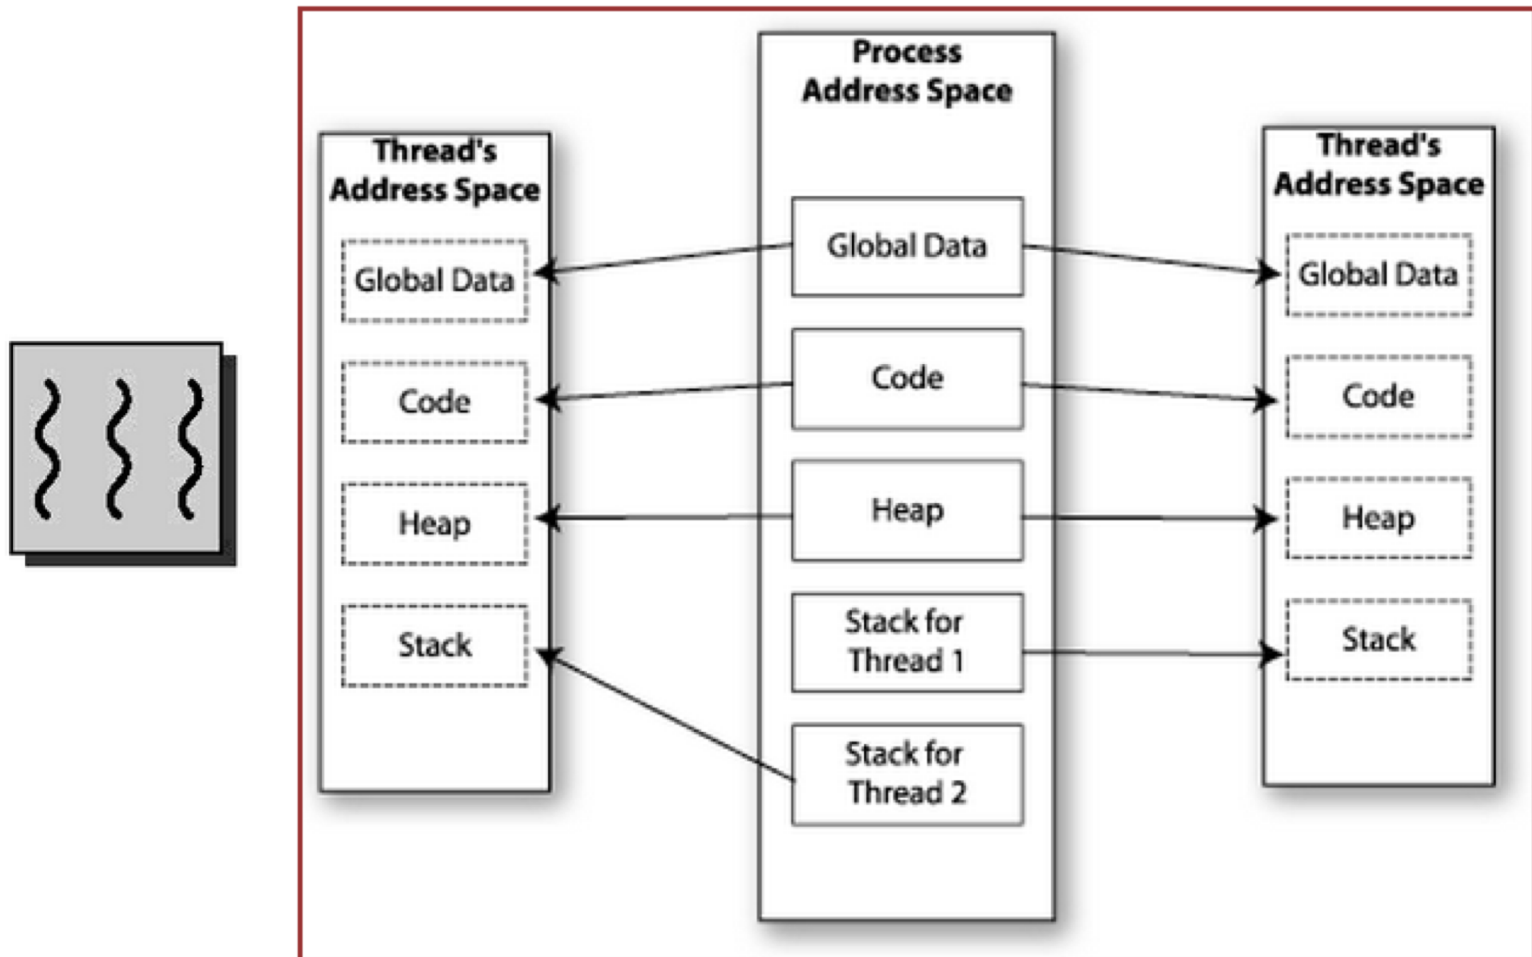
Multiple threads on multiple CPUs



# Memory Layout: Single-Threaded Program



# Memory Layout: Multithreaded Program





# Multitasking

- `$ tr 'abc' 'xyz' | sort -u | comm -23 file1 -`
  - Process 1 (tr)
  - Process 2 (sort)
  - Process 3 (comm)
- Each process has its own address space
- How do these processes communicate?
  - Pipes/System Calls

# Multithreading

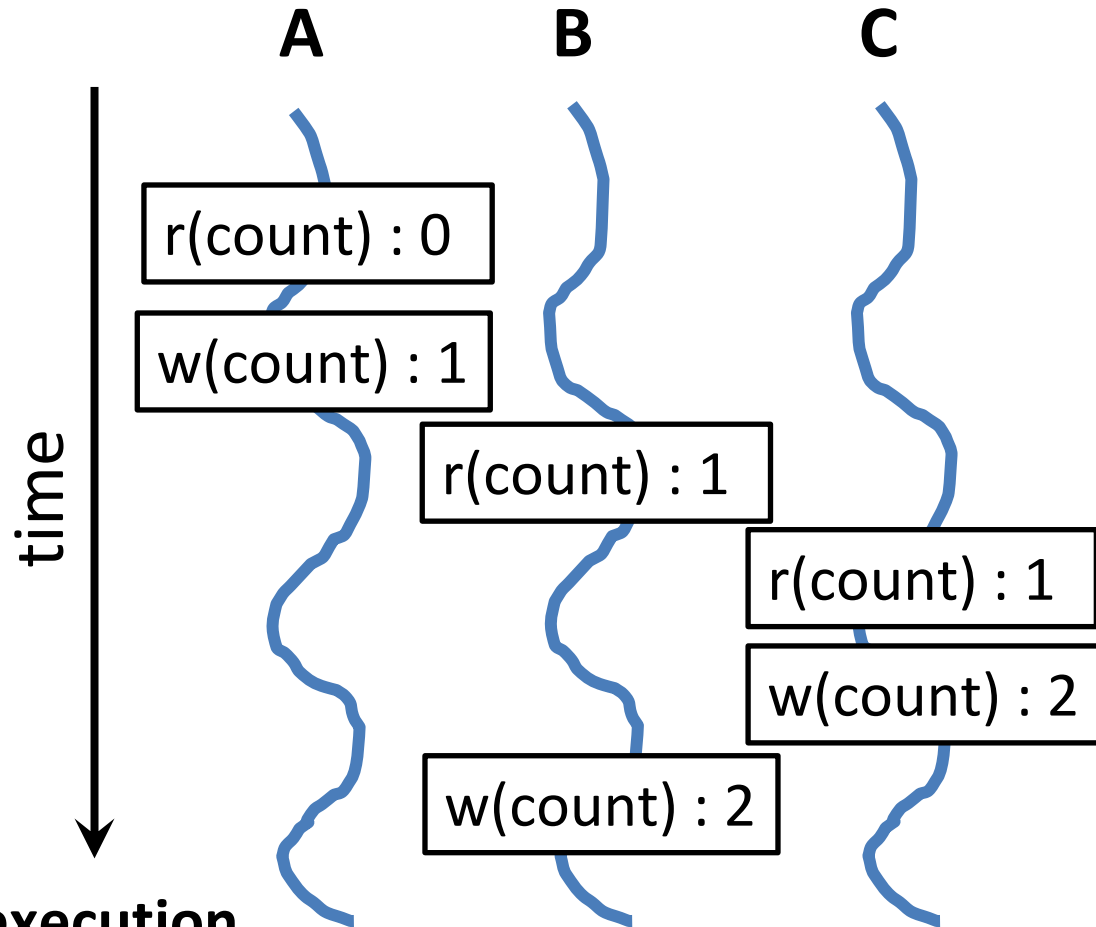
- Threads share all of the process's memory except for their stacks
- => Data sharing requires no extra work (no system calls, pipes, etc.)

# Shared Memory

- Makes multithreaded programming
  - **Powerful**
    - can easily access data and share it among threads
  - **More efficient**
    - No need for system calls when sharing data
    - Thread creation and destruction less expensive than process creation and destruction
  - **Non-trivial**
    - Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)

# Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```



**Result depends on order of execution  
=> Synchronization needed**

# Multithreading & Multitasking: Comparison

- **Multithreading**

- Threads share the same address space
  - Light-weight creation/destruction
  - Easy inter-thread communication
  - An error in one thread can bring down all threads in process

- **Multitasking**

- Processes are insulated from each other
  - Expensive creation/destruction
  - Expensive IPC (interprocess communication)
  - An error in one process cannot bring down another process

# Lab 6

- Evaluate the performance of multithreaded `sort`
- Add `/usr/local/cs/bin` to `PATH`
  - `$ export PATH=/usr/local/cs/bin:$PATH`
- Generate a file containing random single-precision floating point numbers, one per line with no white space
  - `/dev/urandom`: pseudo-random number generator
- Disk quota exceeded
  - <http://www.seasnet.ucla.edu/seasnet-account-quotas/>

# Lab 6

- `od`
  - write the contents of its input files to standard output in a user-specified format
  - Options
    - `-t` : select output format
    - `-N <count>`: Format no more than *count* bytes of input
- `sed, tr`
  - Remove address, delete spaces, add newlines between each float

# Lab 6

- use `time -p` to time the command `sort -g` on the data you generated
- Send output to `/dev/null`
- Run `sort` with the `--parallel` option and the `-g` option: compare by general numeric value
  - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
    - `$ time -p sort -g file_name > /dev/null (1 thread)`
    - `$ time -p sort -g --parallel=[2, 4, or 8] file_name > /dev/null`
  - Record the times and steps in `log.txt`