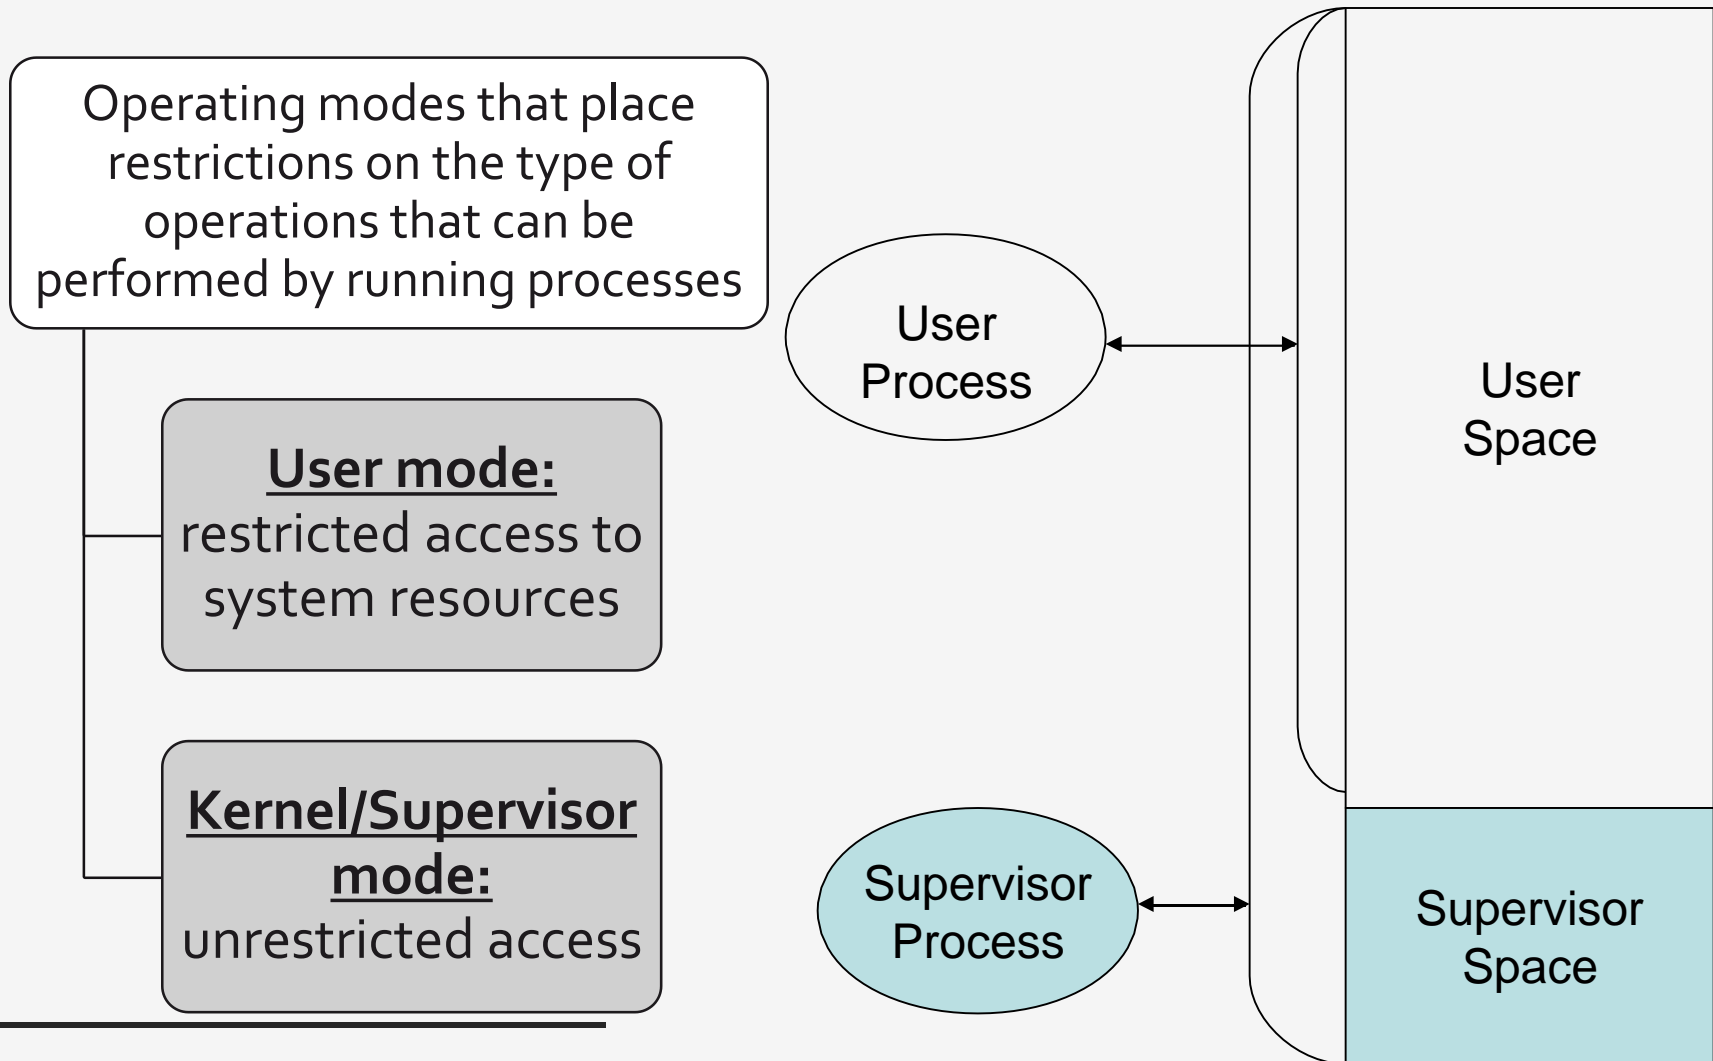


# *CS35L – Winter 2019*

|               |                         |
|---------------|-------------------------|
| Slide set:    | 5.1                     |
| Slide topics: | System call programming |
| Assignment:   | 5                       |



# *Processor Modes*




# *User Mode vs. Kernel Mode*

- Hardware contains a mode-bit, e.g. 0 means kernel mode, 1 means user mode
- User mode
  - CPU **restricted** to unprivileged instructions and a specified area of memory
- Supervisor/kernel mode
  - CPU is **unrestricted**, can use all instructions, access all areas of memory and take over the CPU anytime




# *Why Dual-Mode Operation?*

---

- System resources are shared among processes
  - OS must ensure:
    - **Protection**
      - an incorrect/malicious program cannot cause damage to other processes or the system as a whole
    - **Fairness**
      - Make sure processes have a fair use of devices and the CPU
- 

# *Goals for Protection and Fairness*




## **I/O Protection**

- Prevent processes from performing illegal I/O operations



## **Memory Protection**

- Prevent processes from accessing illegal memory and modifying kernel code and data structures



## **CPU Protection**

- Prevent a process from using the CPU for too long

# *Which Code is Trusted?*

## *The Kernel Only*

- Core of OS software **executing in supervisor state**
- **Trusted software:**
  - Manages hardware resources (CPU, Memory and I/O)
  - Implements protection mechanisms that could not be changed through actions of untrusted software in user space
- **System call interface is a safe way** to expose privileged functionality and services of the processor

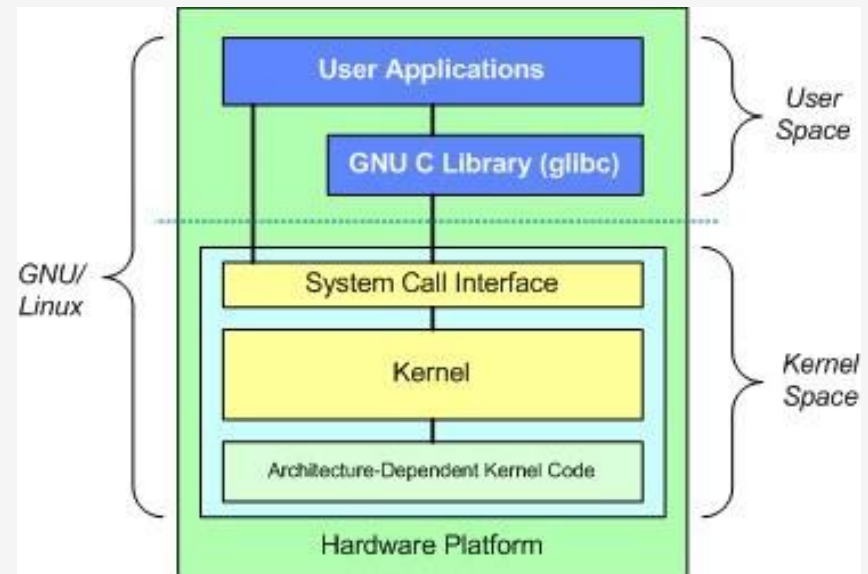



Image by: Tim Jones (IBM)

# *What About User Processes?*


- The kernel executes privileged operations on behalf of untrusted user processes



# *System Calls*

---

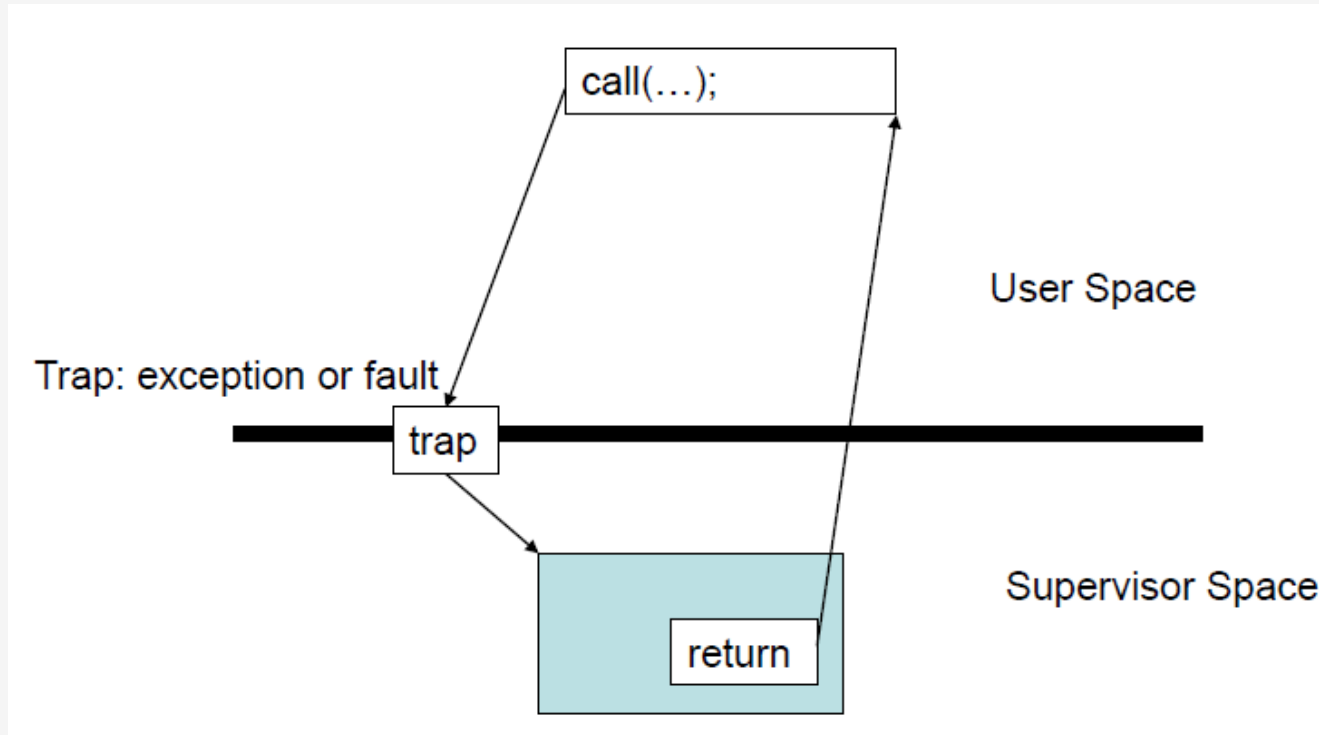
- Special type of function that:
  - Is used by user-level processes to request a service from the kernel
  - Changes the CPU's mode from user mode to kernel mode to enable more capabilities
  - Is part of the kernel of the OS
  - Verifies that the user should be allowed to do the requested action and then does the action (kernel performs the operation on behalf of the user)
  - Is the **only way** a user program can perform privileged operations





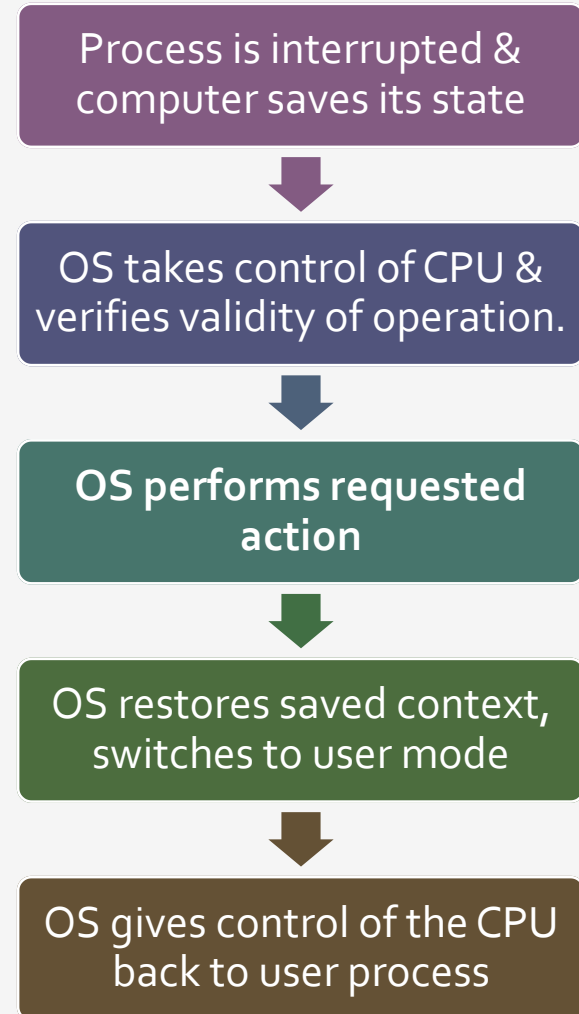
# *System Calls*

- When a system call is made, the program being executed is interrupted and control is passed to the kernel
- If operation is valid the kernel performs it



# *System Call Overhead*

*System calls are expensive and can hurt performance*



# *Example System Calls*

```
#include<unistd.h>
```

- `ssize_t read(int fildes, void *buf, size_t nbyte)`
    - `fildes`: file descriptor
    - `buf`: buffer to write to
    - `nbyte`: number of bytes to read
  - `ssize_t write(int fildes, const void *buf, size_t nbyte);`
    - `fildes`: file descriptor
    - `buf`: buffer to write from
    - `nbyte`: number of bytes to write
  - `int open(const char *pathname, int flags, mode_t mode);`
  - `int close(int fd);`
  - File descriptors
    - 0 `stdin`
    - 1 `stdout`
    - 2 `stderr`
-

# *Example System Calls*

`pid_t getpid(void)`

- Returns the process ID of the calling process

`int dup(int fd)`

- Duplicates a file descriptor `fd`. Returns a second file descriptor that points to the same file table entry as `fd` does.


`int fstat(int filedes, struct stat *buf)`

- Returns information about the file with the descriptor `filedes` into `buf`

```
struct stat {
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t     st_mode;     /* protection */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t  st_blksize;  /* blocksize for file system I/O */
    blkcnt_t   st_blocks;   /* number of 512B blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

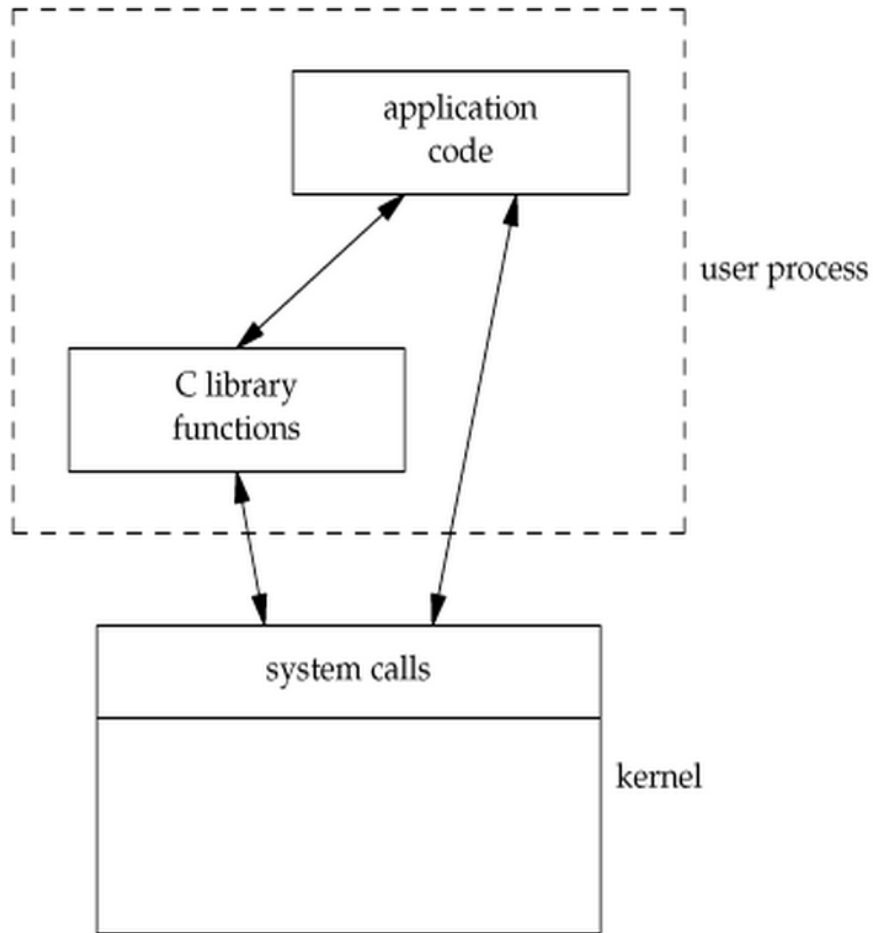
---

# *Library Functions*



- Functions that are a part of standard C library
- To avoid system call overhead use equivalent library functions
  - `getchar, putchar` vs. `read, write` (for standard I/O)
  - `fopen, fclose` vs. `open, close` (for file I/O), etc.
- How do these functions perform privileged operations?
  - They make system calls!

# *So What's the Point?*



Many library functions  
invoke system calls indirectly

So why use library calls?

- Usually equivalent library functions make fewer system calls
- non-frequent switches from user mode to kernel mode => less overhead

# *Unbuffered vs. Buffered I/O*

## Unbuffered

Every byte is read/written by the kernel through a system call

## Buffered

Collect as many bytes as possible (in an internal buffer) and read more than a single byte (into buffer) at a time and use one system call for a block of bytes

*Buffered I/O decreases the number of read/write system calls and the corresponding overhead*