

Week 6

System Call Programming

11 February 2019

CS 35L Lab 4

Jeremy Rotman

Announcements

- Assignment #5 is due Saturday by 11:55pm
- For Assignment #10
 - ◆ Email me to tell me what story you are choosing
 - ◆ [Here is the link to see what stories people have signed up for already](#)
 - Choose a story at least one week before you present
- The Final is on Sunday March 17 3-6pm, and is a common final
 - ◆ Let me know if this creates a schedule conflict ASAP

Outline

- Dual-Mode Operation
- System Calls
- Buffered vs. Unbuffered I/O
- Lab 5

Questions?

Processor Modes

→ The CPU (in Linux) has two distinct modes of operation

◆ Kernel mode

- Unrestricted access
- Can execute any instruction, and reference any memory address
- Assumes it is running trusted software

◆ User mode

- Non-privileged access
- Cannot directly access hardware
- Must use a system call to perform privileged instructions

Why Dual-Mode Operation?

Why Dual-Mode Operation?

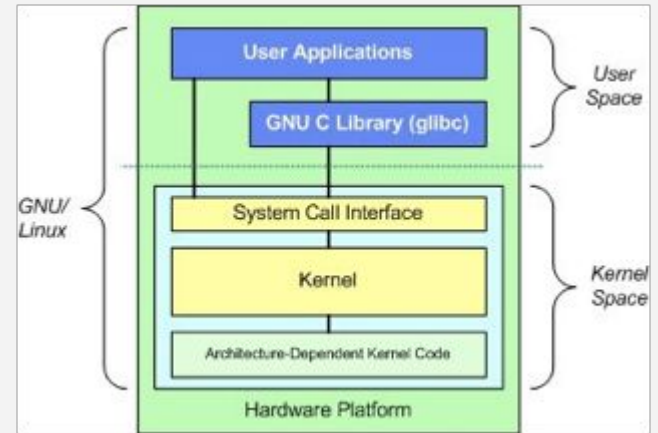
- I/O protection
 - ◆ Input/output cannot be directly controlled by user-created code
- Memory protection
 - ◆ User mode only has access to a set partition of memory
 - ◆ The user is not allowed to access memory addresses that define these bounds, or other addresses outside their partition
- CPU Protection
 - ◆ User mode is not allowed to change things related to the OS's scheduler or timer

Trusted Software

- What is the trusted software that can be run in kernel mode?

Trusted Software

- What is the trusted software that can be run in kernel mode?
 - ◆ Software in the kernel space
 - Cannot be changed from the outside
 - Implements protection mechanisms
- System call interface bridges the gap between User Mode and Kernel Mode
 - ◆ User processes can execute privileged operations through the interface




System Calls

- Used by user-level processes to request service from kernel
- Changes CPU's mode from User to Kernel
- Part of the Kernel of the OS
- Verifies User should be allowed to do operation, then handles operation
- The only way that a user program can perform privileged operations

System Calls

- When a system call is made
- ◆ The program being executed is interrupted
 - ◆ Control is passed to the kernel
 - ◆ If the operation is valid
 - Kernel performs it
 - ◆ Else
 - Throw an exception



System Calls

→ Overhead

System Calls

→ Overhead

- ◆ System calls include a lot of overhead
- ◆ Many things that must be done
 - Process interrupted and computer saves its state
 - OS takes control of CPU and verifies validity of operation
 - OS performs requested action
 - OS restores saved context, switches to user mode
 - OS gives control of the CPU back to user process

Example System Calls

- `ssize_t read(int fildev, void* buf, size_t nbyte)`
 - ◆ *fildev*: file descriptor
 - ◆ *buf*: address of buffer to read into
 - ◆ *nbyte*: maximum number of bytes to read
- `ssize_t write(int fildev, void* buf, size_t nbyte)`
 - ◆ *fildev*: file descriptor
 - ◆ *buf*: address of buffer to read into
 - ◆ *nbyte*: maximum number of bytes to write
- `int open(const char* pathname, int flags, mode_t mode)`
- `int close(int fildev)`
- `fd: 0 = stdin; 1 = stdout; 2 = stderr;`

Example System Calls

- `void exit(int status)`
 - ◆ Terminates process with *status*
- `pid_t fork(void)`
 - ◆ Creates child process
- `pid_t getpid(void)`
 - ◆ Returns the process ID of the calling process
- `int dup(int fildes)`
 - ◆ Duplicates a file descriptor, *fildes*
- `int fstat(int fildes, struct stat* buf)`
 - ◆ Return information about the file with file descriptor *fildes*, into *buf*

C Library Functions

- There are functions that are part of the standard C library that do the same thing
 - ◆ getchar & putchar are similar to read & write (standard I/O)
 - ◆ fopen & fclose are similar to open & close (file I/O)
- These C functions then make system calls of their own
- What is the advantage?

C Library Functions

- There are functions that are part of the standard C library that do the same thing
 - ◆ getchar & putchar are similar to read & write (standard I/O)
 - ◆ fopen & fclose are similar to open & close (file I/O)
- These C functions then make system calls of their own
- What is the advantage?
 - ◆ Library functions make fewer system calls
 - ◆ This reduces the amount of overhead
 - ◆ Efficiency :D

Unbuffered vs. Buffered I/O

→ Unbuffered

- ◆ Every byte is read/written by the kernel through system call

→ Buffered

- ◆ Collect as many bytes as possible (in a buffer) and read/write them all in one system call

→ Buffered I/O reduces overhead because you don't have to switch mode for every byte

- ◆ Thought exercise: when might you want to use unbuffered I/O?
- ◆ Also, you must still always be prepared for buffer overflow attacks

Lab 5

→ Write C code to transliterate bytes

◆ 2 arguments:

- *from*: the bytes to transliterate
- *to*: the bytes to transliterate them into

◆ Works basically the same way as `tr`

- `tr 'abcd' 'nopq' < input_file`
 - Replace 'a' with 'n', 'b' with 'o', etc.

→ You will write 2 programs (one buffered, one unbuffered)

◆ `tr2b`

- Uses `getchar` and `putchar` to read from `stdin` and write to `stdout`

◆ `tr2u`

- Uses `read` and `write` with the `nbyte` argument as 1

Lab 5

- The programs will need to be tested on a very large file
 - ◆ At least 5,000,000 bytes
 - ◆ `head --bytes=num /dev/urandom > output.txt`
- **time** [options] *command* [arguments...]
 - ◆ Output:
 - `real 0m0.145s`
 - Elapsed real time (in [hours:]minutes:seconds)
 - `user 0m0.001s`
 - Total number of CPU-seconds that process spent in user mode
 - `sys 0m0.003s`
 - Total number of CPU-seconds that process spent in kernel mode

Lab 5

- **strace** -o *strace_output command* [arguments...]
- ◆ Intercepts and prints out system calls to stderr or to an output file
 - ◆ Every line will include
 - system call name
 - arguments in parentheses
 - return value

Questions?