

CS35L – *Winter '19*

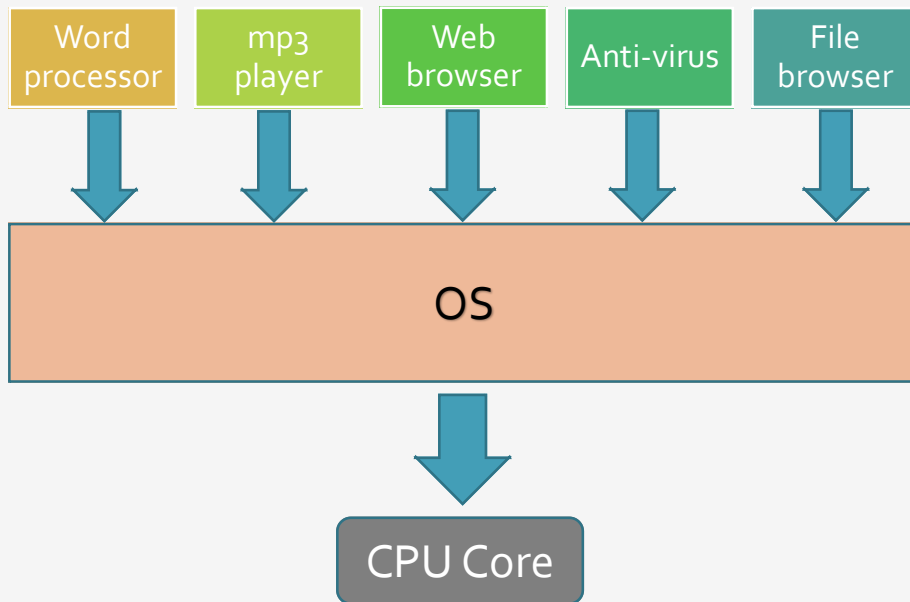
Slide set:	6.1
Slide topics:	Multithreaded performance
Assignment:	6



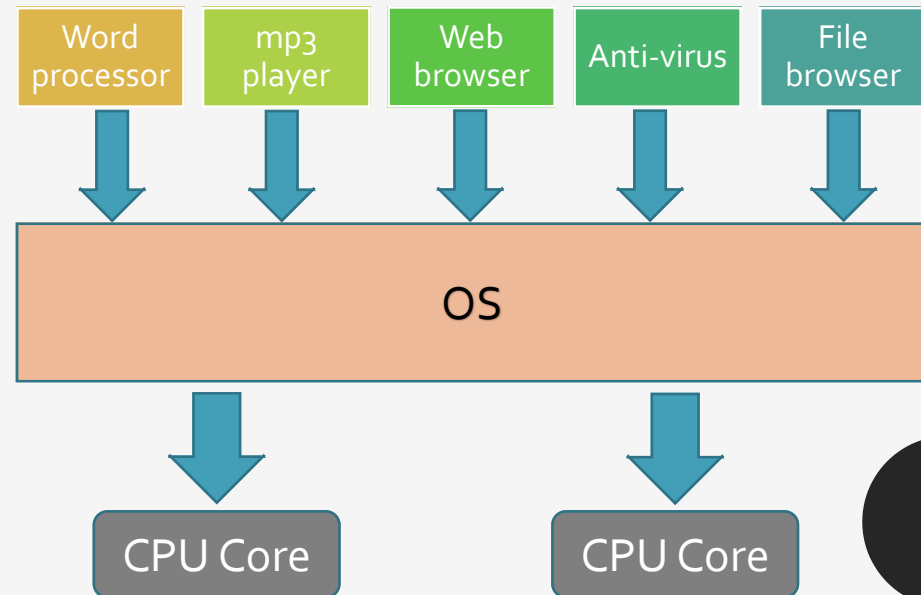
Multiprocessing

The use of multiple CPUs/cores to run multiple tasks simultaneously

Uni-processing system




Multiprocessing system



Parallelism

Executing several computations simultaneously to gain performance



Different types of parallelism

Multitasking

- Several processes are scheduled alternately or possibly simultaneously on a multiprocessing system

Multithreading

- Same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system



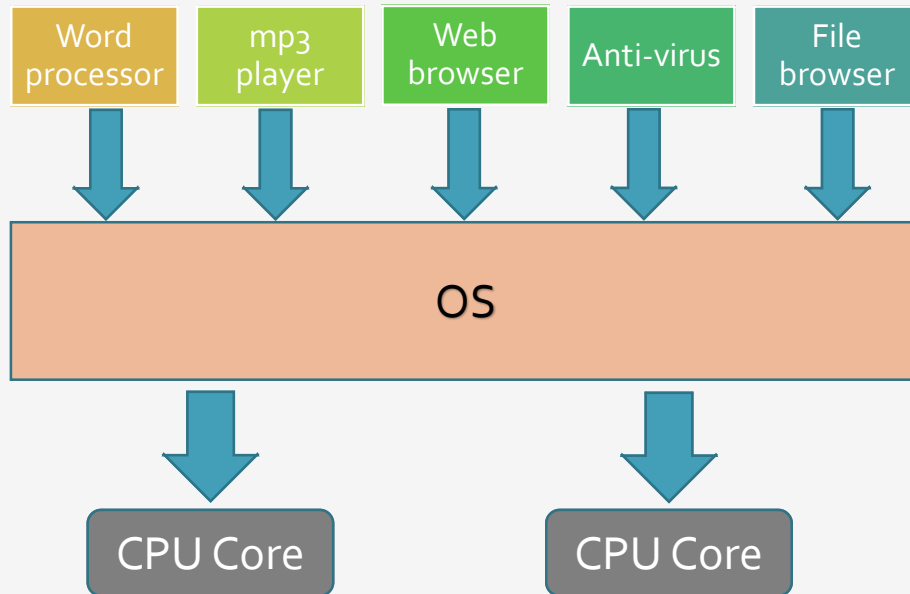
Threads

- A flow of instructions, path of execution within a process
- The smallest unit of processing scheduled by OS
- A process consists of at least one thread
- Multiple threads can be run on:
 - **A uniprocessor (time-sharing)**
 - Processor switches between different threads
 - Parallelism is an illusion
 - **A multiprocessor**
 - Multiple processors or cores run the threads at the same time
 - True parallelism

Multitasking vs. Multithreading

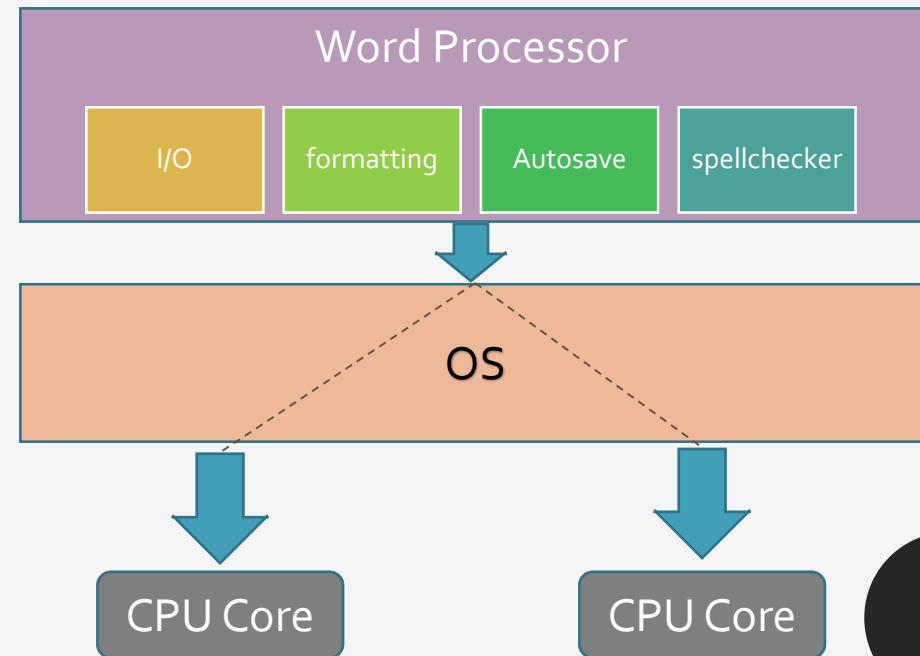
Multitasking

Sharing of computing resources (CPU, memory, devices, etc.) among processes




Multithreading

Sharing of computing resources among threads of a single process.







one process
one thread




one process
multiple threads




multiple processes
one thread per process



multiple processes
multiple threads per process



Multiple threads sharing a single CPU



Multiple threads on multiple CPUs



Multitasking

```
$ tr 'abc' 'xyz' | sort -u |  
comm -23 file1
```

- Process 1 (tr)
- Process 2 (sort)
- Process 3 (comm)

Each process has its own address space

- Area of memory (virtualized) devoted to the process

Communication between processes

- Via Pipes/System Calls

01


Threads share all of the process's memory except for their stacks

02

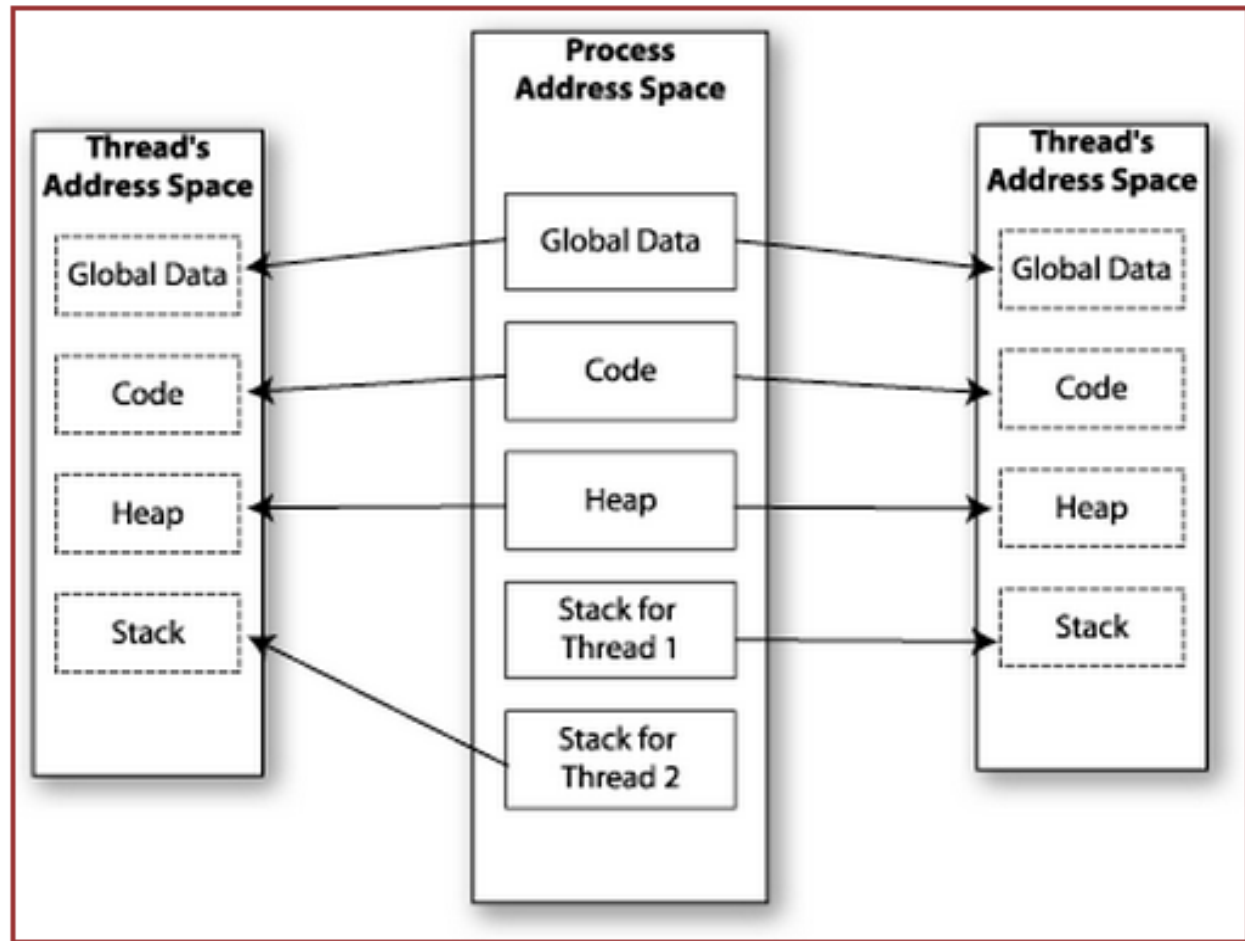
Data sharing requires no extra work (no system calls, pipes, etc.)

Multithreading

Memory Layout: Single-Threaded Program



Memory Layout: Multithreaded Program



- Makes multithreaded programming
 - **Powerful**
 - can easily access data and share it among threads
 - **More efficient**
 - No need for system calls when sharing data
 - Thread creation and destruction less expensive than process creation and destruction
 - **Non-trivial**
 - Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)


Shared Memory

Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```

Result depends on
order of execution

*Synchronization
needed*



Multithreading & Multitasking: A Comparison

Multithreading

- Threads share the same address space
- Light-weight creation/destruction
- Easy inter-thread communication
- An error in one thread can bring down all threads in process

Multitasking

- Processes are insulated from each other
- Expensive creation/destruction
- Expensive IPC (interprocess communication)
- An error in one process cannot bring down another process

- Evaluate the performance of multithreaded sort
- Add /usr/local/cs/bin to PATH
 - `$ export PATH=/usr/local/cs/bin:$PATH`
- Generate a file containing 2^{24} random single-precision floating point numbers, one per line with no white space
 - /dev/urandom: pseudo-random number generator
- Disk quota exceeded
 - <http://www.seasnet.ucla.edu/seasnet-account-quotas/>

Lab 6

- `head`
 - Write the first few lines of a file to output
 - Use `-N` to select number of lines
- `od`
 - Write the contents of its input files to standard output in a user-specified format
 - Options
 - `-t` : select output format
 - `-N <count>`: Format no more than *count* bytes of input
- `sed`, `tr`
 - Remove address, delete spaces, add newlines between each float

Lab 6

Lab 6

- use `time -p` to time the command `sort -g` on the data you generated
- Send output to `/dev/null`
- Run `sort` with the `--parallel` option and the
 - `-g` option: compare by general numeric value
 - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
 - `$ time -p sort -g file_name > /dev/null (1 thread)`
 - `$ time -p sort -g --parallel=[2, 4, or 8] file_name > /dev/null`
 - Record the times and steps in `log.txt`