

The background features abstract, overlapping green geometric shapes in various shades, creating a modern and dynamic look. The shapes are primarily triangular and polygonal, with some areas appearing more translucent than others.

CS 35L

Software Construction Laboratory

Lecture 6.1

12th February, 2019

Logistics

- ▶ Hardware requirement for Week 8
 - ▶ Seeed Studio BeagleBone Green Wireless Development Board
- ▶ Presentations for Assignment 10
 - ▶ https://docs.google.com/spreadsheets/d/1o6r6CKCaB2du3klPflHiquymhBvbn7oP0wkHHMz_q1E/edit?usp=sharing
- ▶ Assignment 5 is due on 16th Feb, 2018 at 11:55pm

Review - Previous Lab

- ▶ System Call Examples
 - ▶ Open, Create, Close
 - ▶ Read, Write
- ▶ Why System Calls?

Assignment 5 - Laboratory

► Review:

- `ch = getchar()`
- `putchar(ch)`
- `int numRead = read(STDIN_FILENO, ch, size)`
- `int numWritten = write(STDOUT_FILENO, ch, size)`

Assignment 5 - Laboratory

- ▶ Write tr2b and tr2u programs in 'C' that transliterates bytes. They take two arguments 'from' and 'to'. The programs will transliterate every byte in 'from' to corresponding byte in 'to'
 - ▶ `./tr2b 'abcd' 'wxyz' < bigfile.txt`
 - ▶ Replace 'a' with 'w', 'b' with 'x', etc
 - ▶ `./tr2b 'mno' 'pqr' < bigfile.txt`
- ▶ tr2b uses `getchar` and `putchar` to read from STDIN and write to STDOUT.
- ▶ tr2u uses `read` and `write` to read and write each byte, instead of using `getchar` and `putchar`. The `nbyte` argument should be 1 so it reads/writes a single byte at a time.
- ▶ Test it on a big file with 5,000,000 bytes
 - ▶ `$ head --bytes=# /dev/urandom > bigfile.txt`

Tr2b.c

- ▶ Write a main function which accepts arguments
 - ▶ `main(int argc, const char* argv[])`
- ▶ Check for the length of arguments
 - ▶ Retrieve first argument in `char * from`, second argument in `char * to`
 - ▶ Compare the lengths of `from` and `to`; If not same, throw an error and exit
 - ▶ You can use `strlen` to get lengths
- ▶ To throw an error, write to `stderr`
- ▶ To exit, write `exit(1)`
- ▶ Check if 'from' has duplicates
 - ▶ In a loop, take input from `stdin` (till you reach eof of `stdin`) using `getchar()`
 - ▶ Check if the character you just retrieved is a part of `from`; if yes then put the corresponding character in `stdout` with `putchar()`

Tr2u.c

- ▶ Repeat the same procedure as in tr2b.c except replace:
 - ▶ getchar() with read
 - ▶ putchar() with write

Time and strace

- ▶ **time [options] command [arguments...]**
- ▶ **Output:**
 - ▶ -real 0m4.866s: elapsed time as read from a wall clock
 - ▶ -user 0m0.001s: the CPU time used by your process
 - ▶ -sys 0m0.021s: the CPU time used by the system on behalf of your process
- ▶ **strace: intercepts and prints out system calls.**
- ▶ `-$ strace -c ./tr2b 'AB' 'XY' < input.txt`
- ▶ `-$ strace -c ./tr2u 'AB' 'XY' < input.txt`

Additional Information

- ▶ www.cs.uregina.ca/Links/class-info/330/SystemCall_IO/SystemCall_IO.html
- ▶ courses.engr.illinois.edu/cs241/sp2009/Lectures/04-syscalls.pdf
- ▶ www.bottomupcs.com/system_calls.shtml

Assignment 5 - Homework

- ▶ Rewrite sfrob using system calls (sfrobu)
- ▶ sfrobu should behave like sfrob except:
 - ▶ If stdin is a regular file, it should initially allocate enough memory to hold all data in the file all at once
 - ▶ It outputs a line with the number of comparisons performed
- ▶ Functions you'll need: *read*, *write*, and *fstat* (read the man pages)
- ▶ Measure differences in performance between sfrob and sfrobu using the *time* command
- ▶ Estimate the number of comparisons as a function of the number of input lines provided to sfrobu

Assignment 5 - Homework

- ▶ Write a shell script “sfrobs” that uses tr and the sort utility to perform the same overall operation as sfrobu (support -f option as well)
- ▶ Use pipelines (do not create temporary files)
- ▶ Encrypted input -> tr (decrypt) -> sort (sort decrypted text) -> tr (encrypt) -> encrypted output

Assignment 5 - Homework

- ▶ Measure any differences in performance between sfrob and sfrobu using the time command.
- ▶ Run your program on inputs of varying numbers of input lines, and estimate the number of comparisons as a function of the number of input lines
- ▶ Use the time command to compare the overall performance of sfrob, sfrobu, sfrobs, sfrobu -f and sfrobs -f

Assignment 5 - Homework

- ▶ Run your program on inputs of varying numbers of input lines, and estimate the number of comparisons as a function of the number of input lines.
- ▶ Varying number of input lines => number of words
- ▶ Number of comparisons => keep a counter in the frobcmp() function to check how many times it is being called

Assignment 5 - Homework

- ▶ Refer to *Read, Write, Open, Close* System Calls
- ▶ Reserved File Descriptors
 - ▶ 0 - stdin
 - ▶ 1 - stdout
 - ▶ 2 - stderr
- ▶ `int fstat(int fd, struct stat *buf)`
 - ▶ Returns information about the file with the descriptor fd into buf

Parallelism and Multi Threading

Types of Resources

▶ CPU

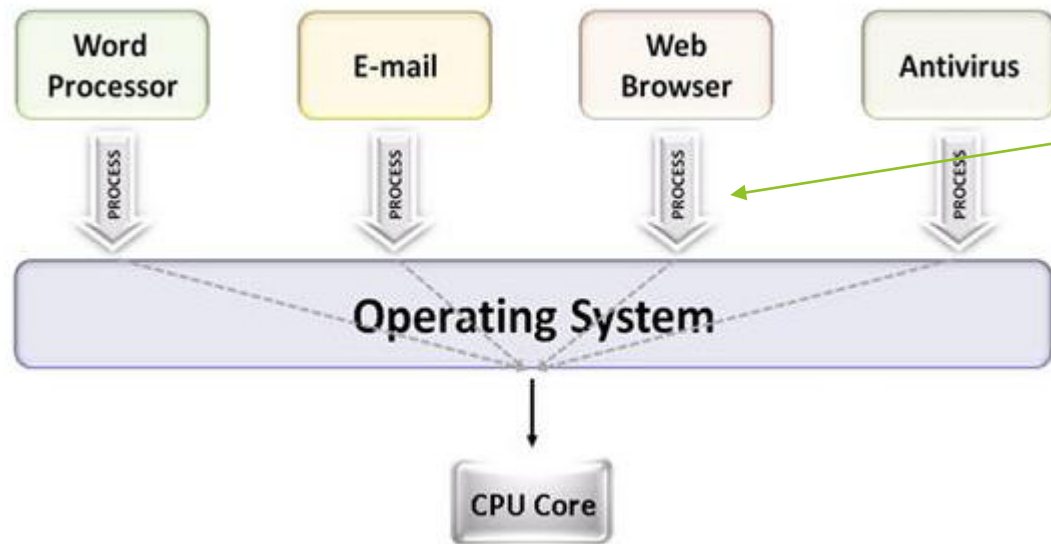
- ▶ It is an active resource
- ▶ Can be used by only one runtime entity
- ▶ Can be multiplexed in time (time sharing)

▶ Memory

- ▶ Passive resource
- ▶ Can be shared among multiple runtime entities
- ▶ Can be multiplexed in space (allocated)

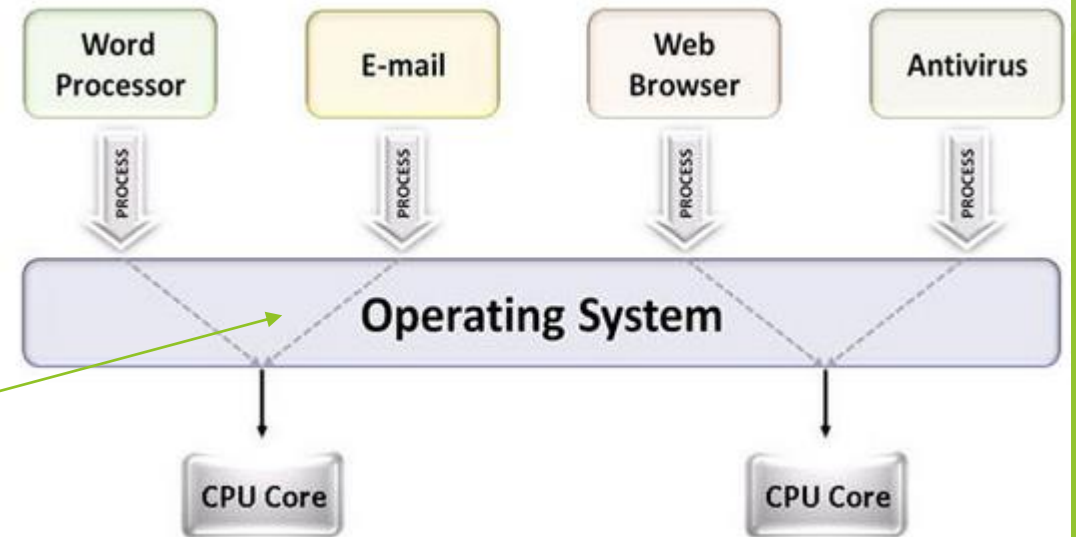
Multiprocessing

- The use of multiple CPUs/cores to run multiple tasks simultaneously



**Multiprocessing
system**

**Uni-processing
system**



Parallelism

- ▶ Executing several computations simultaneously to gain performance
- ▶ Different forms of parallelism
- ▶ **Multitasking**
 - ▶ Several processes are scheduled alternately or possibly simultaneously on a multiprocessing system
- ▶ **Multithreading**
 - ▶ Same job is broken logically into pieces (threads) which may be executed simultaneously on a multiprocessing system

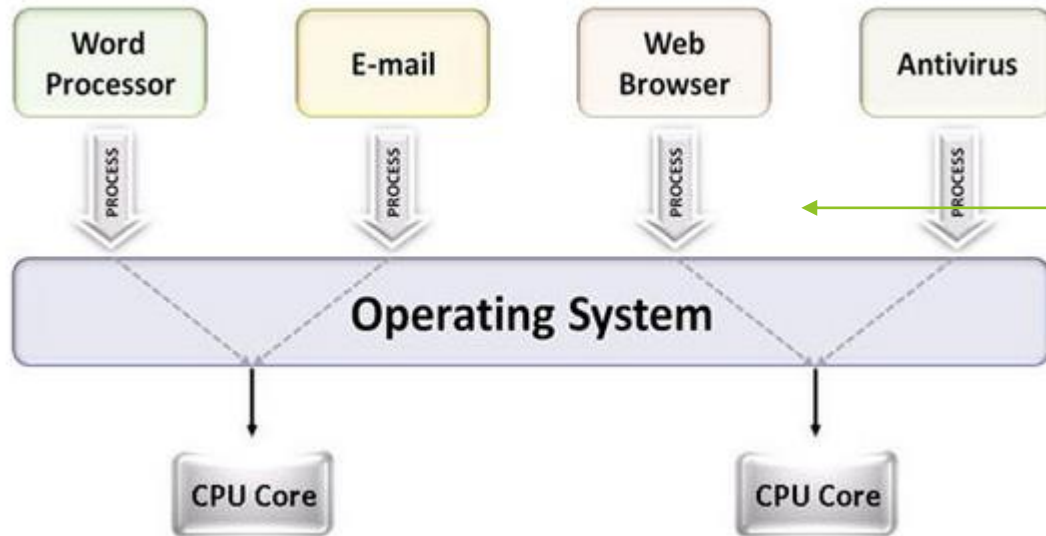
What is a Thread?

- ▶ A flow of instructions, path of execution within a process
- ▶ The smallest unit of processing scheduled by OS
- ▶ A process consists of at least one thread
- ▶ Multiple threads can be run on:
 - ▶ A uniprocessor (time-sharing)
 - ▶ Processor switches between different threads
 - ▶ Parallelism is an illusion
- ▶ A multiprocessor
 - ▶ Multiple processors or cores run the threads at the same time
 - ▶ True parallelism

Process vs Threads

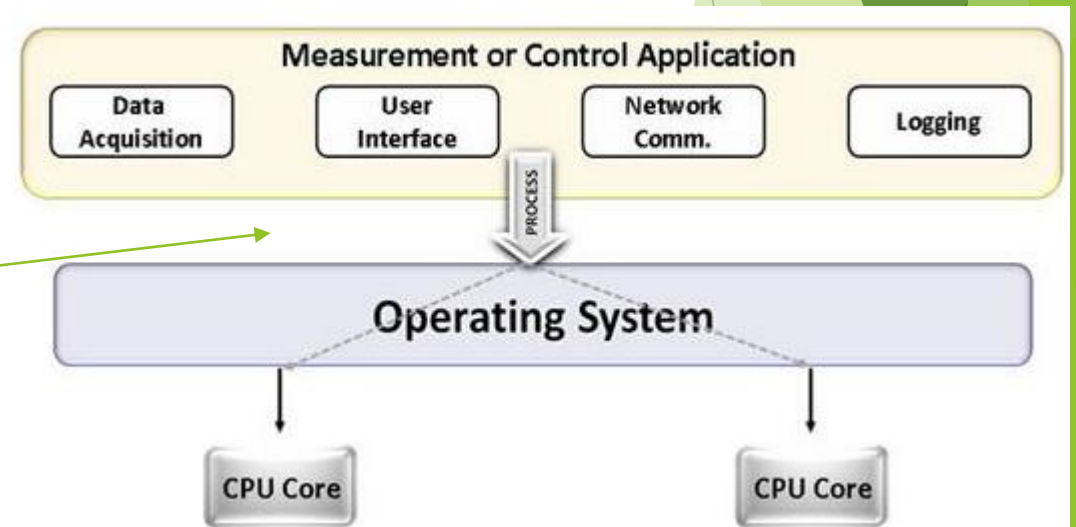
- ▶ Different processes see separate address spaces
 - ▶ good for protection, bad for sharing
- ▶ All threads in the same process share the same memory (except stack)
 - ▶ good for sharing, bad for protection
 - ▶ each thread can access the data of other thread

Multitasking vs. Multithreading




Multitasking

Multithreading



Multithreading




Multithreading & Multitasking: Comparison

- ▶ Multithreading
 - ▶ Threads share the same address space
 - ▶ Light-weight creation/destruction
 - ▶ Easy inter-thread communication
 - ▶ An error in one thread can bring down all threads in process
- ▶ Multitasking
 - ▶ Processes are insulated from each other
 - ▶ Expensive creation/destruction
 - ▶ Expensive Inter Process Communication
 - ▶ An error in one process cannot bring down another process

Multithreading & Multitasking

Multiple threads sharing a single CPU



Multiple threads on multiple CPUs



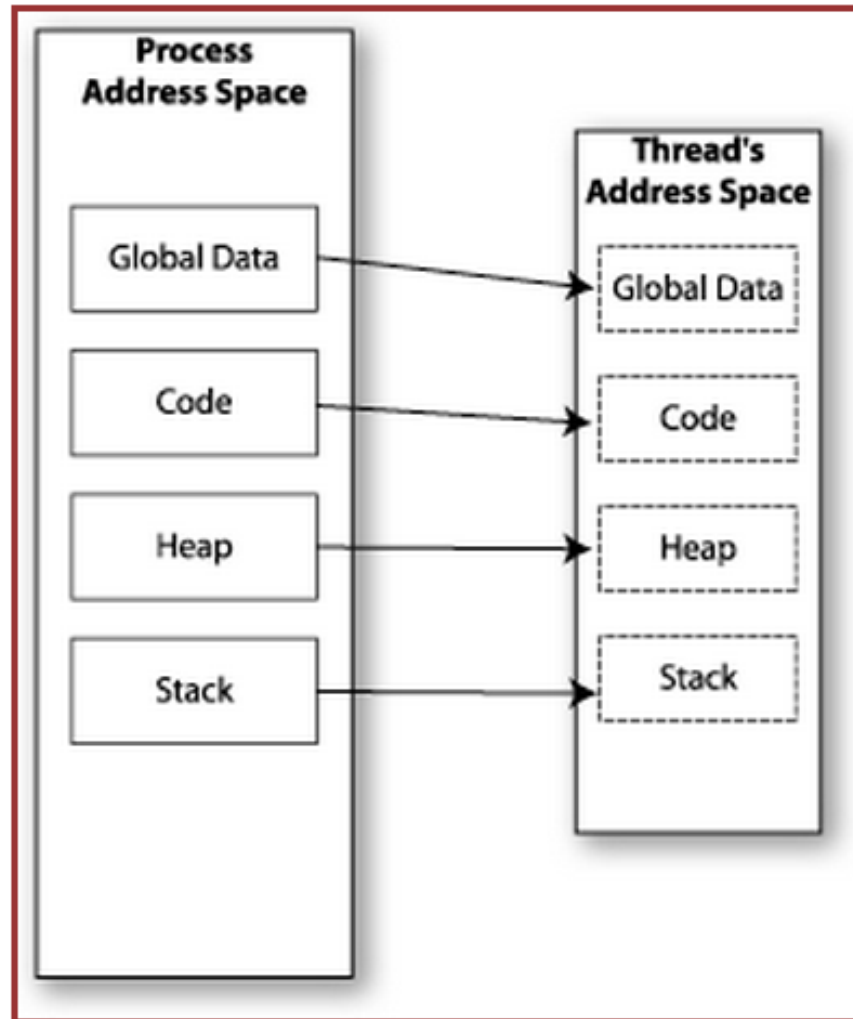

Multitasking

- ▶ `$ tr 'abc' 'xyz' | sort -u | comm -23 file1 -`
 - ▶ Process 1 (tr)
 - ▶ Process 2 (sort)
 - ▶ Process 3 (comm)
- ▶ Each process has its own address space
- ▶ How do these processes communicate?
 - ▶ Pipes/System Calls

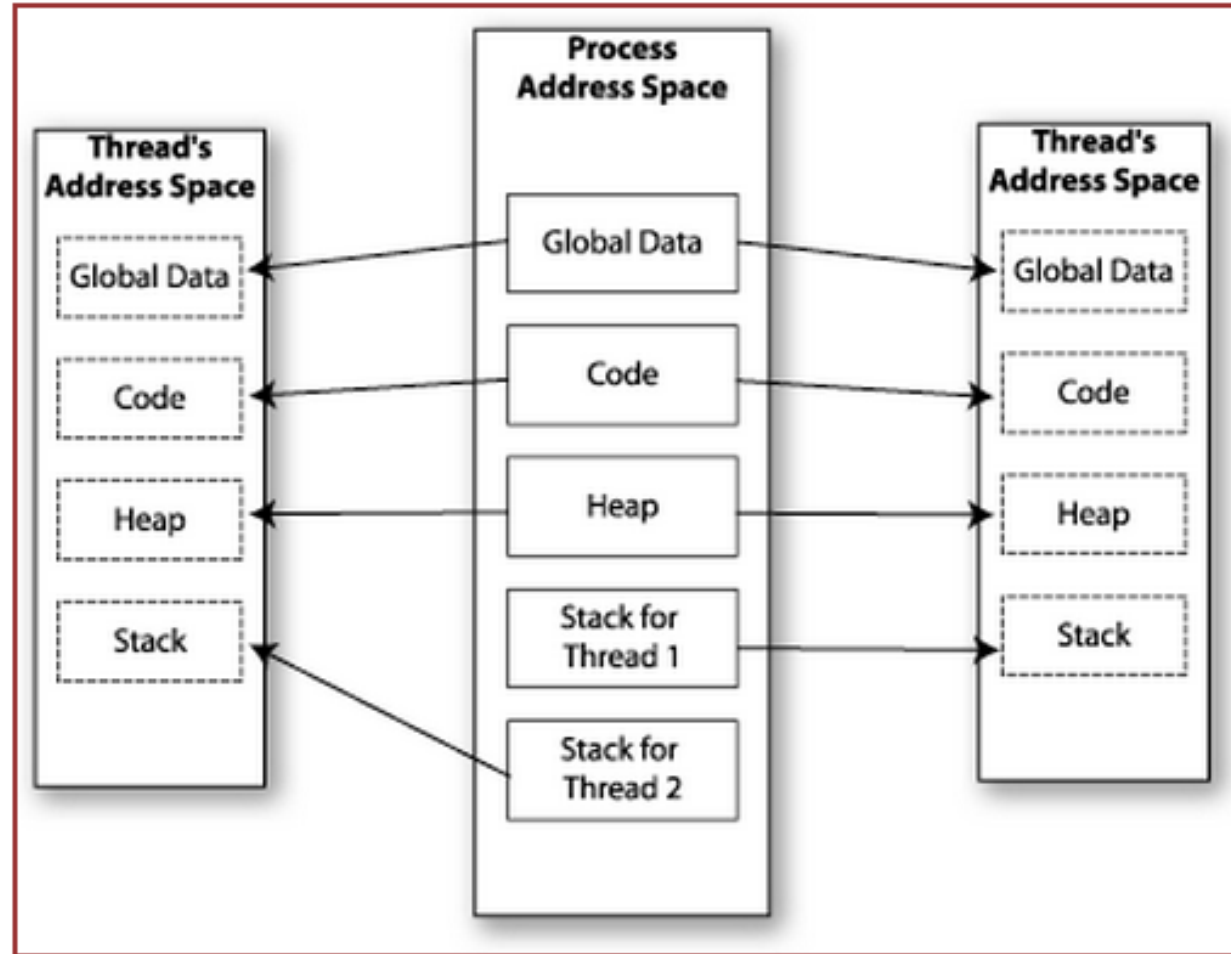
Multithreading

- ▶ Threads share all of the process's memory except for their stacks
- ▶ Data sharing requires no extra work (no system calls, pipes, etc.)

Memory Layout: Single-Threaded Program



Memory Layout - Multi Threaded Program



Shared Memory

- ▶ Makes multithreaded programming
- ▶ Powerful
 - ▶ can easily access data and share it among threads
- ▶ More efficient
 - ▶ No need for system calls when sharing data
 - ▶ Thread creation and destruction less expensive than process creation and destruction
- ▶ Non-trivial
 - ▶ Have to prevent several threads from accessing and changing the same shared data at the same time (synchronization)


Process/Thread Synchronization

- ▶ Why is it needed?
- ▶ Because threads share the same resources, we need synchronization
- ▶ To prevent inconsistency

Race Condition

```
int count = 0;  
void increment()  
{  
    count = count + 1;  
}
```

Result depends on order of execution
=> Synchronization needed



Presentations

- ▶ **Today's Presentation:**

- ▶ Richard Cheng
- ▶ Ethan Kwan

- ▶ **Next up:**

- ▶ Niyant Narang
- ▶ Rio Sonoyama

Questions?