
CS 35L- Software Construction Laboratory

Winter 19

TA: Guangyu Zhou

Course Information

- Presentation Schedule

Sena Ji	N/A	IBM Research Develops Fingernail Sensor to Monitor Disease Progression	Seiji Otsu	N/A	Button Offers Instant Gratification for Those Plagued by Airplane Noise	Yingge Zhou	Siyan Dong	Facial recognition
---------	-----	---	------------	-----	--	-------------	------------	--------------------


- No Class Next Monday Feb 18 (President day holiday)

System call programming and debugging

Week 6

Review: Processor Modes

- To understand system calls, first we need to distinguish between **supervisor (kernel) mode** and **user mode** of a CPU
 - **Mode** bit used to distinguish between execution on behalf of OS & behalf of user.
- Modern operating system supports these two modes
 - Supervisor (Kernel) mode: processor executes every instruction in it's hardware repertoire
 - User mode: can only use a subset of instructions



Review: Kernel vs User Mode

Kernel Mode

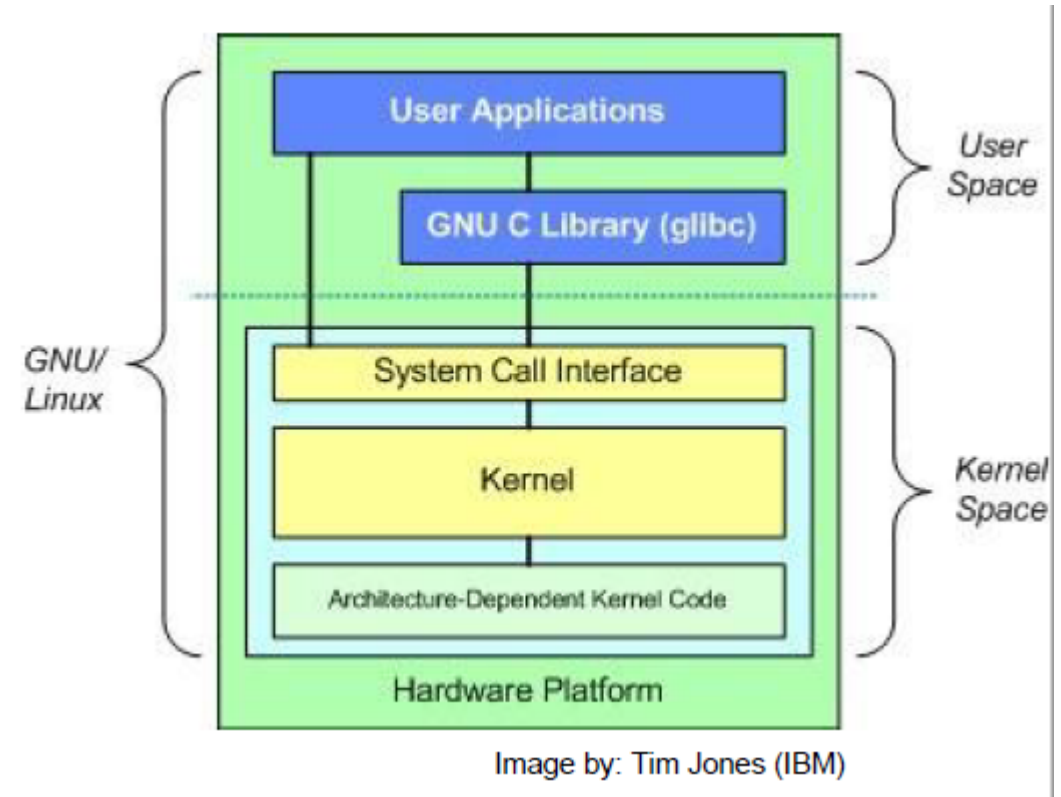
- When CPU is in **kernel mode**, the code being executed can access any memory address and any hardware resource.
- Hence kernel mode is a very privileged and powerful mode.
- If a program crashes in kernel mode, the entire system will be halted.

User Mode

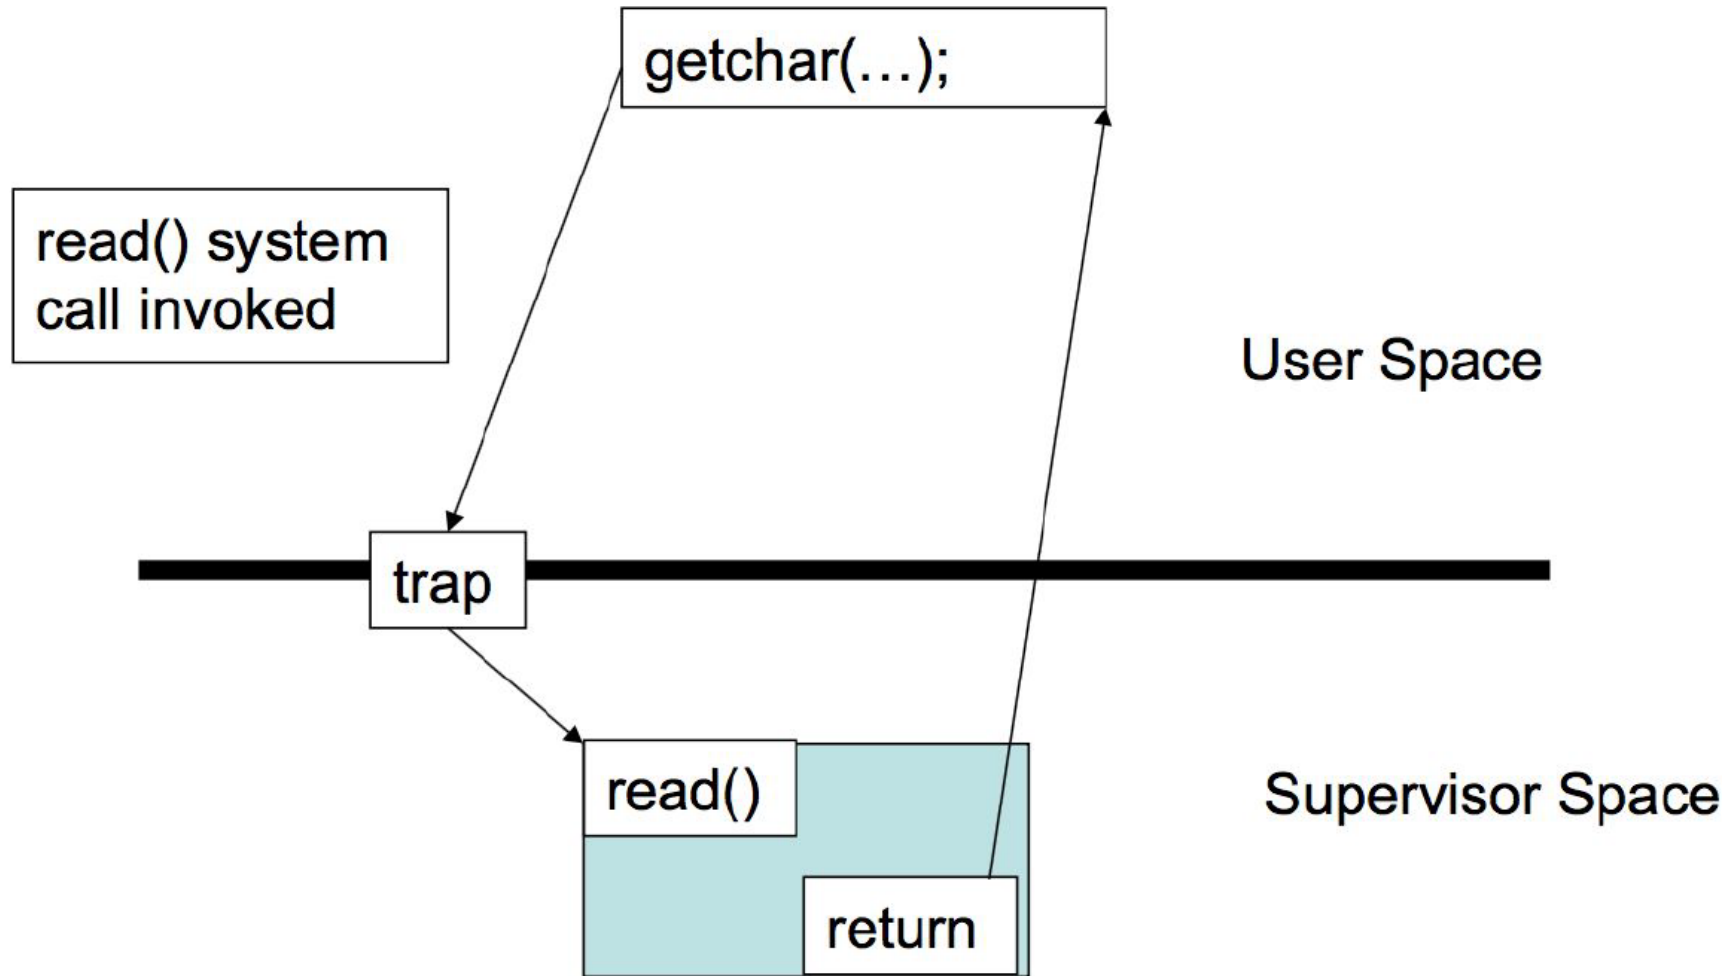
- When CPU is in **user mode**, the programs don't have direct access to memory and hardware resources.
 - In user mode, if any program crashes, only that particular program is halted.
 - That means the system will be in a safe state even if a program in user mode crashes.
 - Hence, most programs in an OS run in user mode.
-

Review: The Kernel

- Code of the OS **executes** in **supervisor** state
- Trusted software
 - manages hardware resources (CPU, memory, and I/O)
 - Implements protection mechanisms that could not be changed through actions of untrusted software in user space
- **System call interface** is a safe way to expose privileged functionality



Reminder of how System calls work



Trap: System call causes a switch from user mode to kernel mode

Reminder of how System calls work

- 1. program get to the system call in the user's code

```
int res = sys_call(a_few_params)
```

- 2. puts the parameters on the stack
- 3. performs a system 'trap' -- hardware switch

now in system mode


- 4. operating system code may copy large data structures into system memory
 - 5. starts operation...
-

Reminder of how System calls work (cont.)

- 6. operation complete!
 - 7. if necessary copies result data structures back to user program's memory
return to user mode
 - 8. user program puts return code into res(the return value from the system call)
 - 9. program recommences
-

More detailed steps for making a system call:

- 1-3: push parameters on stack
- 4: invoke the system call
- 5: put code for system call on register
- 6: trap to the kernel
- 7-10: since a number is associated with each system call, system call interface invokes/dispatch intended system call in OS kernel and return status of the system call and any return value
- 11: increment stack pointer



System calls and Library calls usage

- Library calls
 - executed in the **user program**
 - may perform **several task**
 - may call system calls
 - System calls
 - executed by **the operating system**
 - perform simple **single** operations
-

Types of system calls

System calls can be roughly grouped into five major categories:

1. Process Control
 - load
 - execute
 - end, abort
 - create process (for example, fork on Unix-like systems)
 - terminate process
 - get/set process attributes
 - wait for time, wait event, signal event
 - allocate, free memory
 2. File Management
 - create file, delete file
 - open, close
 - read, write, reposition
 - get/set file attributes
-

Types of system calls

3. Device Management

- request device, release device
- read, write, reposition
- get/set device attributes
- logically attach or detach devices

4. Information Maintenance

- get/set time or date
- get/set system data
- get/set process, file, or device attributes

5. Communication

- create, delete communication connection
 - send, receive messages
 - transfer status information
 - attach or detach remote devices
-

System calls—Examples I

- `ssize_t read(int fildes, void *buf, size_t nbyte)`
 - `fildes`: file descriptor
 - `buf`: buffer to write to
 - `nbyte`: number of bytes to read
 - `ssize_t write(int fildes, const void *buf, size_t nbyte)`
 - `fildes`: file descriptor
 - `buf`: buffer to write to
 - `nbyte`: number of bytes to write
 - `int open(const char *pathname, int flags, mode_t mode)`
 - `int close(int fd)`
-

System calls—Examples II

- `pid_t getpid (void)`
 - returns the process id of the calling process
 - `int dup(int fd)`
 - Duplicates a file descriptor `fd`. Returns a second file descriptor that points to the same file table entry as `fd` does.
 - `int fstat(int fildes, struct stat *buf)`
 - Returns information about the file with the descriptor `fildes` to `buf`
 - *Why are these system calls and not just regular library functions?*
-

System calls vs library call conventions

- Library functions often return pointers
 - `FILE *fp = fopen("cs35l","r")`
 - return NULL for failure
 - System calls usually return an integer
 - `int res=system_call_function(a_few_args)`
 - Where the return value
 - `res >= 0` → all is well
 - `res < 0` → failure
 - See the global variable `errno` for more info
-

More hints for Assignment 5

- **Lab 5: Programs tr2b and tr2u in 'C':**

- Take two arguments 'from' and 'to'.
- Transliterate every **byte** in 'from' to corresponding byte in 'to'
- Report an error **from** and **to** are not the same length, or if **from** has duplicate bytes

```
./tr2b 'abc' < bigfile.txt => Error
```

```
./tr2b 'abc' 'wxyz' < bigfile.txt => Error
```

```
./tr2b 'abad' 'wxyz' < bigfile.txt => Error
```

- **tr2b: getchar and putchar**

- **tr2u: read and write**

- `ret = read(STDIN_FILENO, buf, nbyte=1);`
- `ret = write(STDOUT_FILENO, buf, nbyte=1);`

The following symbolic constants shall be defined for file streams:

STDIN_FILENO File number of stdin; 0.

STDOUT_FILENO File number of stdout; 1.

STDERR_FILENO File number of stderr; 2.

Homework 5

- Recall Homework 4!
 - Rewrite *sfrob* using **system calls** (*sfrobu*)
 - *sfrobu* should behave like *sfrob* except
 - If stdin is a regular file, it should initially allocate enough memory to **hold all data in the file all at once**
 - It outputs a line with the number of comparisons performed
 - System call functions you'll need: **read, write, and fstat**
-

Homework 5

- Measure differences in **performance** between *sfrob* and *sfrobu* using the ***time*** command
 - Estimate the number of **comparisons** as a function of the number of input lines provided to *sfrobu*
 - Write a shell script “*sfrobs*” that uses *tr* and the *sort* utility to perform the same overall operation as *sfrob*
 - Encrypted input -> *tr* (decrypt) -> *sort* (sort decrypted text) -> *tr* (encrypt) -> encrypted output
-

More hints for Assignment 5

- **Homework 5: Encrypted sort revisited**
- `int fstat(int filedes, struct stat *buf)`
 - Returns information about the file with the descriptor `filedes` into `buf`
- `sfrobu` should behave like `sfrob` except:
 - If `stdin` is a regular file, it should initially allocate enough memory to hold all data in the file all at once

/ File Information */*

ret = fstat(STDIN_FILENO, &buf);

fileSize = buf.st_size + 1;

/ Setup Initial Buffer */*

input = (char) malloc(sizeof(char) * fileSize);*

```
struct stat {
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t     st_mode;     /* protection */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t  st_blksize;  /* blocksize for file system I/O */
    blkcnt_t   st_blocks;   /* number of 512B blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

};