

# Week 7

# Multithreading

20 February 2019

CS 35L Lab 4

Jeremy Rotman

# Announcements

- Assignment #6 is due Saturday by 11:55pm
- For Assignment #10
  - ◆ **Email me to tell me what story you are choosing**
  - ◆ [Here is the link to see what stories people have signed up for already](#)
    - Choose a story at least one week before you present
    - Try to avoid political topics

# Outline

- Multithreading and Multitasking
- POSIX threads
- Assignment 6

Questions?

# Multiprocessing


What is Multiprocessing?

# Multiprocessing

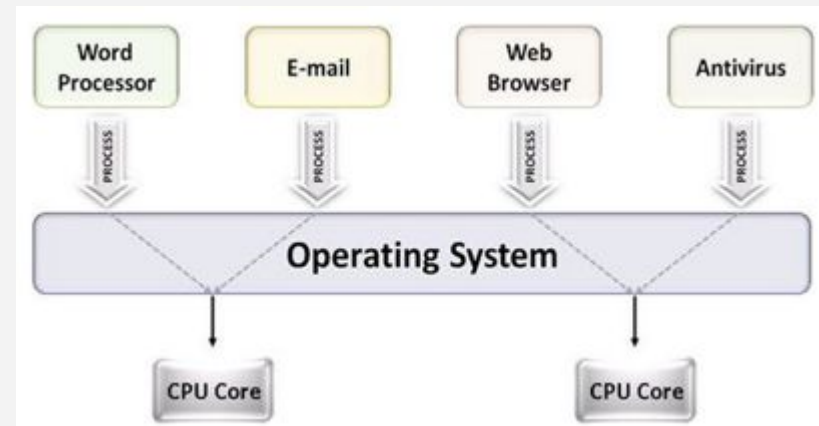
## What is Multiprocessing?

The use of multiple CPUs/cores to run multiple tasks simultaneously

### Uniprocessing System



### Multiprocessing System



# Parallelism

# Parallelism

- Executing several computations simultaneously to gain performance



# Parallelism

- Executing several computations simultaneously to gain performance
- Two primary forms of parallelism
  - ◆ Multitasking
    - Several processes are scheduled in an alternating pattern
    - Potentially, a simultaneous pattern if on a multiprocessing system
  - ◆ Multithreading
    - The same job is split into logical pieces (threads)
    - Threads may be executed simultaneously in a multiprocessing system

# Threads

- A flow of instructions, or path of execution within a process
- Smallest unit of process scheduled by the OS
- A process consists of at least one thread

# Multithreading

→ Multiple threads can be run on:

- ◆ A Uniprocessor

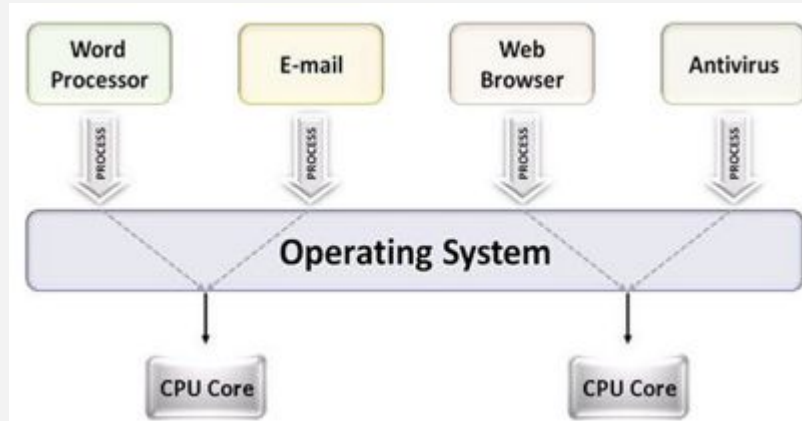
- Processor switches between threads
- An illusion of parallelism

- ◆ A Multiprocessor


- Multiple processes or cores run the threads at the same time
- True parallelism

# Multitasking vs. Multithreading

## Multitasking



## Multithreading



# Using Both

Multithreading

No

Yes

Multitasking

No



One process  
One thread



One process  
Multiple threads

Yes




Multiple processes  
One thread each



Multiple processes  
Multiple threads in each

# Scheduling Threads

Multiple threads sharing a single CPU

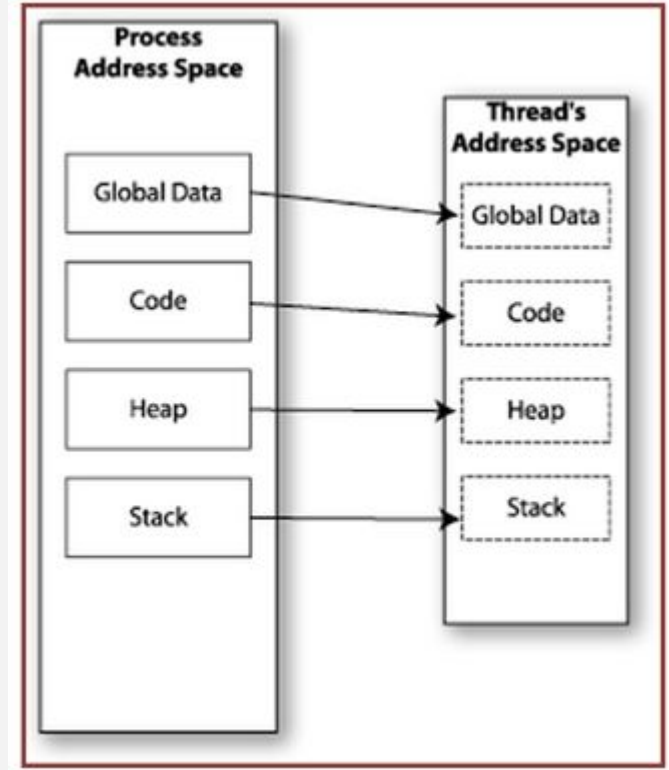


Multiple threads on multiple CPUs



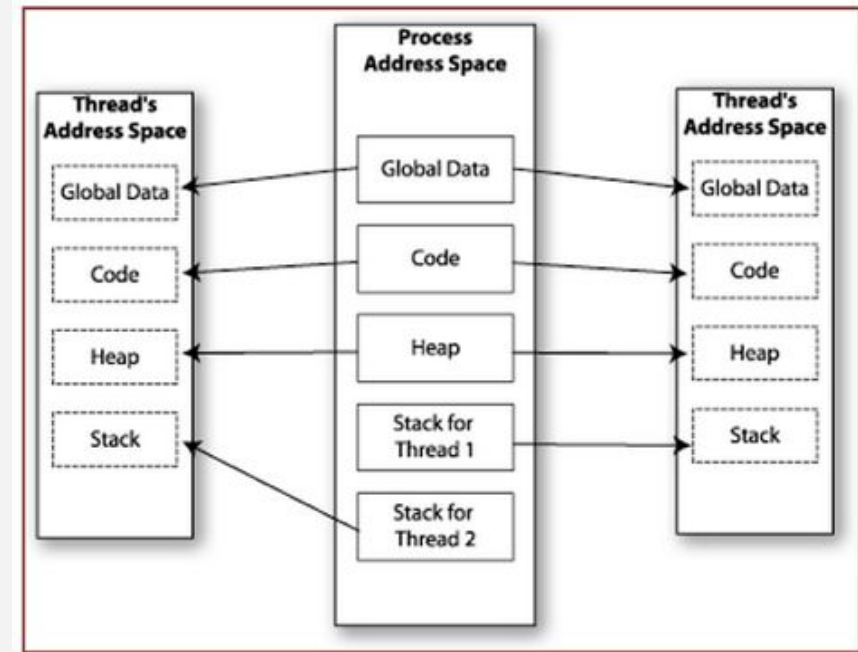
# Memory Layout

- Single-threaded program
- The process allocates space for the thread



# Memory Layout

- Multi-threaded program
- The process allocates individual space for each thread's stack
- Each thread shares the rest of the process's space






# Shared Memory

- Every process receives its own memory
- Every thread receives part of its parent process's memory
  - ◆ Threads from the same parent process share memory
- Data sharing requires no extra work
  - ◆ Processes must interact with other processes with system calls, pipes, etc.
- Makes Multithreading efficient
  - ◆ Also less expensive
- Also makes Multithreading difficult
  - ◆ Race Conditions

# Race Conditions

→ Threads are often scheduled in small chunks




→ What happens if two threads both rely on specific values of data in memory?

# Race Conditions

```
i = 0;  
while(i < 10) {  
    i++;  
    cout << i;  
}
```


Multiple  
threads  
sharing a  
single CPU



# Race Conditions

```
i = 0;  
while(i < 10) {  
    i++;  
    cout << i;  
}
```

Multiple  
threads  
sharing a  
single CPU



# Race Conditions


Output:

3

3

3

Multiple  
threads  
sharing a  
single CPU



# Multithreading and Multitasking

## → Multithreading

- ◆ Threads share the same address space
  - Lightweight creation/destruction
  - Easy inter-thread communication
  - An error in one thread can bring all of the threads down

## → Multitasking

- ◆ Processes are insulated from each other
  - Expensive creation/destruction
  - Expensive IPC
  - An error in one process cannot bring down another process

# POSIX Threads (pthreads)

Built in functions to create and manage threads:

- `pthread_create`
  - ◆ Create a new thread within a process
- `pthread_join`
  - ◆ Waits for another thread to terminate
- `pthread_equal`
  - ◆ Compare thread IDs to check if they are equal
- `pthread_self`
  - ◆ Return ID of the calling thread
- `pthread_exit`
  - ◆ Terminates the currently running thread

# pthread\_create

- Create a new thread and make it executable
- Can be called any number of times from anywhere in the code
- Return
  - ◆ Success: 0
  - ◆ Failure: Error Code



# pthread\_create

```
int pthread_create(pthread_t* tid, const pthread_attr_t*  
attr, void* (myfunction)(void*), void* arg);
```

- tid: unique identifier for new thread
- attr: object that holds attributes
  - ◆ NULL for default
- my\_function: function the thread will execute
- arg: a *single* argument that can be passed to my\_function
  - ◆ NULL for no argument

# pthread\_join

- Makes originating thread wait for the completion of all of its spawned reads
- Without join, the originating thread would exit as soon as it's done with the join
  - ◆ A spawned thread can be aborted even if it is in the middle of its chore
- Return
  - ◆ Success: 0
  - ◆ Failure: Error Code

## pthread\_join

```
int pthread_join(pthread_t tid, void** status);
```

- tid: ID of thread to wait on
- status: Exit status of target thread is stored in the location pointed to by \*status
  - ◆ NULL if no status needed

The job to be done by  
the thread, one void\*  
argument

Loop to create threads,  
thread IDs are just ints

Loop to wait for each  
thread to close

```
#include <stdio.h>
#include <stdlib.h>
#include <assert.h>
#include <pthread.h>
#include <unistd.h>

#define NUM_THREADS 5

void *perform_work(void *arguments){
    int index = *((int *)arguments);
    int sleep_time = 1 + rand() % NUM_THREADS;
    printf("THREAD %d: Started.\n", index);
    printf("THREAD %d: Will be sleeping for %d seconds.\n", index, sleep_time);
    sleep(sleep_time);
    printf("THREAD %d: Ended.\n", index);
}

int main(void) {
    pthread_t threads[NUM_THREADS];
    int thread_args[NUM_THREADS];
    int i;
    int result_code;

    //create all threads one by one
    for (i = 0; i < NUM_THREADS; i++) {
        printf("IN MAIN: Creating thread %d.\n", i);
        threads[i] = i;
        thread_args[i] = i;
        result_code = pthread_create(&threads[i], NULL, perform_work, &thread_args[i]);
        assert(!result_code);
    }

    printf("IN MAIN: All threads are created.\n");

    //wait for each thread to complete
    for (i = 0; i < NUM_THREADS; i++) {
        result_code = pthread_join(threads[i], NULL);
        assert(!result_code);
        printf("IN MAIN: Thread %d has ended.\n", i);
    }

    printf("MAIN program has ended.\n");
    return 0;
}
```

## Lab 6

- Evaluate multithreaded sort
  - ◆ Make sure to use the path `/usr/local/cs/bin`
- Generate a file with  $2^{24}$  random floating point numbers
  - ◆ This may show up as  $2^{2^4}$  in the spec on your laptop
- Get this by reading bytes from `/dev/random` as single-precision floats with **od**
  - ◆ `-t`: selects output format
  - ◆ `-N count`: formats no more than *count* bytes
- Use **sed** and **tr** to format this to be one number per line
  - ◆ Also remove whitespace

# Homework 6

## → Ray Tracing

- ◆ Mimics propagation of light through objects
- ◆ Simulates effects of a single light ray as it's reflected or absorbed by objects in the images
- ◆ You'll be given code that does this

## → Run the original file

## → Then, make it multithreaded

## → Run the multithreaded version and compare the images

# Homework 6

## → Modifying the Program

- ◆ Include `<pthread.c>` in `main.c`
- ◆ Use `pthread_create` and `pthread_join` in `main.c`
- ◆ Link with `-lpthread` flag (LDLIBS target)

## → Make clean check

- ◆ Outputs “1-test.ppm”
- ◆ To see “1-test.ppm”
  - Ubuntu: `sudo apt-get install gimp`
  - Inxsr: X Forwarding
  - `gimp 1-test.ppm`
  - Alternatively, transfer to your own computer

Questions?